

Room correction by inversion in REW — a step-by-step guide

giacomo

2026-09-15

Contents

Room correction by inversion in REW	3
About this document	3
Contents	4
Part I — What inversion actually does	4
1. The method in one paragraph	4
2. What it can correct, and what it provably cannot	5
3. Do you still need X801? Yes — and here is the proof	5
The answer, stated plainly	6
4. The one law you cannot escape	6
5. The failure this guide exists to prevent	7
Part II — Decisions to take before you measure	7
6. One microphone position, or several?	7
Why it works, and why it works especially well here	7
The caveats, honestly	7
How far apart? Measure it, do not guess	8
20 cm is smaller than most guides recommend, on purpose	8
Size the cluster from how the listener actually sits	9
Recommended	9
If you stay single-point	9
7. With several positions, do you still need the FDW?	10
Measured: the two really are independent	10
8. How much regularisation — choosing the FDW cycles	11
The three constraints that pin N down	11
The floor is about modes. It does not apply to a cancellation — use $N = 8$	13
What 12 cycles does to the feature that caused the trouble	13
Two things the FDW is not	13
9. Which smoothing, and why the smoothing menu will not help you	13
If you ever do need smoothing to stick	14
“Smooth the spatial average instead of windowing the captures”	14
10. The correction band	15
Why not push the top edge higher	15
11. Below 80 Hz, correct the sum — not each channel	16
The failure this section exists to prevent	16
The one-line reason	17
The fix is the divisor, not the target	17
Where 80 Hz comes from	17
Building it: two divisions, spliced at 80 Hz	18
Both factors stay cut-only — nothing has to be allowed to boost	18
The equivalent form, and why it is the worse way to build it	18
What it delivers	19
This is not a local invention	19
Building L+R when only separate L and R sweeps exist	19
12. What a magnitude filter cannot do — and the one cheap way round it	20
The failure	20

The half of it that is free	20
What it costs, and what it makes the filter	20
This makes the filter mixed phase — but the harmless kind	21
What this is, in the wider landscape	21
When to reach for it	21

Part III — The procedure 21

TL;DR — the eleven steps, briefly	21
Step 1 — Measure the five positions	24
[!] How low to start the sweep — and why not 10 Hz	24
[!] One reference speaker for the whole set — this is what makes 3c possible	25
Step 2 — Set the window on every original capture, before averaging	25
[!] X801 must not be in the measurement list yet	25
[!] The order is load-bearing: window first, average second	26
L and R showing different Ref Times is correct, not a fault	26
[!] The window is reversible. The traces built from it are not.	27
Step 3 — Reduce the captures to three divisors	27
[!] Cross corr align has no role in this procedure	28
[!] Never equalise L against R	30
The 6.0206 dB is exact, and it has been verified twice	31
“Never average L with R” — and why LR-SP is not a violation of it	32
[!] Multiply, <i>then</i> take minimum phase — never the other order	32
[!] Never window X801	33
Step 4 — Minimum phase, first time	33
[!] The LF tail is not optional — this guide said “no” and was wrong	34
Step 5 — Build the target	35
The check that matters: does the band edge join what is above it?	36
Frequency is the only knob that matters — not slope, not “match the speaker”	36
Why this should clear two of the three acceptance-test failures	37
[!] Ask for more lift than the room offers and you get no correction at all	37
More bass comes from a deeper cut, not a taller shelf	38
The delivered scoop runs deeper than the drawn one — draw shallower than you want to hear	38
If the scoop keeps landing thin: don’t use one	39
Step 6 — Stop and verify	39
Step 7 — The division	40
If the 20–200 Hz group-delay gate fails near the low edge	40
What max gain 0 dB actually does: the filter only ever cuts	41
[!] The feature that caused 1.3 seconds of ringing is 3.5 dB deep	42
Step 8 — Minimum phase, second time	42
This is where the LF-tail error hurts most	42
Step 9 — Bake the crossover correction in, last	43
Step 10 — Export	43
Step 11 — Accept or reject, before deploying	44
[!] Do not verify a five-position filter at one position	44
Step 12 — Prepare the export bundle for open-media-drc	44
[!] L.txt/R.txt/LR.txt must be a sweep — not the traces this	45
procedure has been building all along	45

Part IV — Reference 46

R1. The IR window dialog, field by field	46
Frequency resolution = $1 / (\text{right width})$. Settled by experiment.	46
Route B — a fixed window instead of the FDW	46
R2. Choosing a window shape	47
[!] Textbook window figures do not apply to a half window	47
Why a time window is a smoothing kernel	47
R3. Minimum phase — the two places you take it, and the one place you must not	49
The phase ledger	49
R4. How to invert: $\div$, $1/A$ and $1/ A $	49
R5. The two boost guards in V5.40 beta	50
R6. Getting the room’s decay figure	50
R7. X801 — what it is, and do not redesign it	51

[!] X801.rephase on disk does not match X801.wav	52
Not to be confused with Xo801.wav	52
R8. Acceptance tests	52
[!] The first test used to be in FFT bins. That was wrong	52
Seeing all three tests in REW, and comparing two filters	53
Thresholds, so you can judge audibility rather than pass/fail	54
What the guards cost	54
How the gated-tone test works	54
[!] A single tone length is not a measurement	55
Worked example: R8 blinded by processing, then not	55
R9. Does the source correlation invalidate §11?	55
Correlation, in one paragraph	55
What the room measurement assumes	56
What is actually on the discs	56
And yet it changes nothing — because of the clamp	56
The general lesson, which is the reason this appendix exists	56
R10. Troubleshooting	57
R11. Glossary	57

Part V — Automating the procedure 58

13. What it does, and does not, do for you	58
14. Input measurements, and the config file	58
Optional session cleanup	59
15. Steps 2–10, one API call each	60
Optional refinement PEQs: 82 Hz and 530 Hz	60
16. Building the target	61
17. Output files	62
18. Rebuilding at a different FDW, or after a tweak	62
19. Quirks the API did not document	62
R12. Appendix — how the REW API was explored	63

Room correction by inversion in REW

A step-by-step guide, with the reasoning behind each choice.

Procedure revision: [rew-inversion-v2026.08.18](#)

REW interface baseline: V5.40 beta 131. The measurements used for the new multi-position example were made with V5.40 beta 122; the action names below are those in beta 131.

V5.40 replaced the old Actions panels with actions in the **All SPL graph right-click menu**. It also replaced the division operation’s obsolete **Regularisation** percentage with an optional **Max gain** control. If the screen in front of you has a **Regularisation** field, it is not the beta UI documented here.

About this document

This is a **procedure**, not a lab notebook. Every number in it has been measured, but the measurements, the false starts and the corrections live in The working history is not repeated here; this document is the current procedural reference.

Scope. Producing a pair of FIR room-correction filters for a stereo pair, in REW, by **inverting the measured response** — as opposed to fitting parametric filters with Auto EQ. The filters are minimum phase, exported as WAV, and convolved by BruteFIR.

The configuration assumed throughout is DRC-120.blue, which is the current one:

Loudspeakers	B&W Nautilus 801 , 3-way, crossovers at 350 Hz and 4 kHz, third order
Placement	120 cm from the front wall; microphone 449 cm from the front wall
Room	≈ 58 m ³ , irregular (7.40 m long, 4.186 → 5.986 m wide, ceiling 2.4 → 3.0 m, open corridor)

Schroeder frequency	$\approx 166 \text{ Hz}$
Correction band	20 – 225 Hz
Sample rate	48 kHz throughout
Crossover correction	X801.wav, built in rePhase, 131072 taps
Convolution engine	BruteFIR
Data	inversion example: ../DRC-120.blue/120.blue.mdat; multi-position examples: ../DRC-120.blue/120.blue.screens.multimeas.mdat and ../DRC-120.blue/120.blue.multi.pt.mdat

Not covered: subwoofers, speaker placement, Auto EQ, room treatment.

Contents

- [Part I — What inversion actually does](#)
 - [1. The method in one paragraph](#)
 - [2. What it can correct, and what it provably cannot](#)
 - [3. Do you still need X801?](#)
 - [4. The one law you cannot escape](#)
 - [5. The failure this guide exists to prevent](#)
 - [Part II — Decisions to take before you measure](#)
 - [6. One microphone position, or several?](#)
 - [7. With several positions, do you still need the FDW?](#)
 - [8. How much regularisation — choosing the FDW cycles](#)
 - [9. Which smoothing, and why the smoothing menu will not help you](#)
 - [10. The correction band](#)
 - [11. Below 80 Hz, correct the sum — not each channel](#)
 - [12. What a magnitude filter cannot do — and the one cheap way round it](#)
 - [Part III — The procedure](#)
 - [Part IV — Reference](#)
 - [R1. The IR window dialog, field by field](#)
 - [R2. Choosing a window shape](#)
 - [R3. Minimum phase — the two places you take it](#)
 - [R4. How to invert: \$\div\$, \$1/A\$ and \$1/|A|\$](#)
 - [R5. The two boost guards in V5.40 beta](#)
 - [R6. Getting the room's decay figure](#)
 - [R7. X801 — what it is, and do not redesign it](#)
 - [R8. Acceptance tests](#)
 - [R9. Does the source correlation invalidate §11?](#)
 - [R10. Troubleshooting](#)
 - [R11. Glossary](#)
 - [Part V — Automating the procedure](#)
 - [13. What it does, and does not, do for you](#)
 - [14. Input measurements, and the config file](#)
 - [15. Steps 2–10, one API call each](#)
 - [16. Building the target](#)
 - [17. Output files](#)
 - [18. Rebuilding at a different FDW, or after a tweak](#)
 - [19. Quirks the API did not document](#)
 - [R12. Appendix — how the REW API was explored](#)
-

Part I — What inversion actually does

1. The method in one paragraph

Measure the loudspeaker in the room. Decide what you *want* that response to be — the **target**. Divide the target by the measurement. The quotient is a correction curve: wherever the room was 6 dB loud the quotient is

−6 dB, wherever it was 4 dB shy the quotient is +4 dB. Turn that curve into an impulse response and convolve it with the music. Because convolution in time is multiplication in frequency, **measurement** × **correction** = **target** — the room is cancelled by construction.

That is the whole idea, and it is exact. Everything difficult about the method comes from one place: **the division is an ill-posed inverse**. Where the measurement is very small, the quotient is very large; where the measurement has a razor-thin notch, the quotient has a razor-thin resonance. An unguarded inversion faithfully reproduces every defect of the measurement, including the ones that are not properties of the room at all.

So the whole craft of this method is deciding what the inversion is allowed to see. Every setting in this guide is a form of that decision.

2. What it can correct, and what it provably cannot

Any measured response factors — uniquely — into two parts:

$$H(\text{measured}) = H(\text{minimum phase}) \times H(\text{excess phase})$$

- **H(min)** carries all the magnitude. Its phase is not free: it is fixed by the magnitude through the Hilbert relation. Know the magnitude and you know the phase.
- **H(excess)** has magnitude exactly 1 everywhere. It is pure phase — pure delay and rotation, invisible in an amplitude plot.

This procedure produces a **minimum-phase** filter (that is what steps 4 and 8 enforce). A minimum-phase filter is derived from a magnitude. Therefore:

A minimum-phase inversion corrects H(min) completely and H(excess) not at all — not “poorly”, but exactly zero.

That is not a limitation to be engineered around; it is the point. Look at what **H(excess)** contains here:

component of the excess phase	invert it?	why
bulk propagation delay	no	it is the speed of sound; inaudible, and inverting it is acausal
room reflections	no	position-specific, changes over 10 cm, and its inverse is acausal — it pre-rings
crossover all-pass	yes	a property of the loudspeaker; position-independent; deterministic; known in closed form

Two of the three must be left alone. The one that should be corrected is the one the minimum-phase route cannot reach. **That is the entire job of X801**, and it is why the crossover correction is a separate, model-based filter rather than something the inversion discovers.

Restricting the inversion to minimum phase is therefore a choice, not a shortcoming. A **complex** inversion — dividing by the raw measured response, phase and all — would correct all three rows of that table, including the two that must not be corrected. It would invert the room’s reflections, which needs an acausal filter, pre-rings, and is valid only at the microphone point. Steps 4 and 8 exist to make sure that never happens. The price of that safety is the third row, and X801 is what pays it.

3. Do you still need X801? Yes — and here is the proof

The question is fair: if inversion corrects the response, shouldn’t it correct the phase everywhere too, by definition?

It would — if it were a **complex** inversion. It is not: steps 4 and 8 force the filter to minimum phase, and §2 gives the reason — a complex inversion would also invert the room’s reflections, which must not be inverted. And the crossover’s contribution is precisely the kind of thing a minimum-phase inversion is blind to:

A third-order Butterworth crossover sums to

$$\frac{1}{(s+1)(s^2+s+1)} + \frac{s^3}{(s+1)(s^2+s+1)} = \frac{s^2 - s + 1}{s^2 + s + 1}$$

— a **second-order all-pass**. Magnitude 1.000000 across the whole band, 359° of phase rotation. A 4th-order Linkwitz-Riley sums to the same form with Q = 0.707 instead of 1.0. **Any crossover that sums to flat magnitude puts its phase where magnitude-derived correction cannot see it.**

Measured on the actual file, over the whole band 5 Hz – 23 kHz:

|X801.wav| min -0.00000 dB max +0.00000 dB

And the decisive test — take the minimum-phase version of X801.wav:

	result
peak	1.000000 at sample 0
energy anywhere but sample 0	2.96×10^{-14} %

The minimum-phase copy of X801 is a unit impulse. Not approximately — to fourteen decimal places. Ask the inversion to correct the crossover and it correctly answers “there is nothing here”.

And it earns its place. Cascaded onto the measured channels (best-fit pure delay removed, so what is left is the phase distortion that actually smears transients), X801 cuts total excess-phase deviation over 200 Hz–15 kHz by ~**30 %** — L 674° → 469° peak-to-peak, R 695° → 466°. The correction is concentrated at the **woofer/midrange crossover, 200–700 Hz**, where R’s excess-phase spread drops from 67° to 24° and about **1.3 ms of group delay** is flattened. Above ~1 kHz it straightens a gentle phase tilt — visible on a wrapped-phase overlay, perceptually close to a pure delay — and does not touch the fine ripple or the tweeter crossover. The audible part is the lower midrange.

The answer, stated plainly

Yes, you still need X801, and no amount of inversion will ever replace it. They address disjoint parts of the response — magnitude and magnitude-implied phase on one side, crossover all-pass on the other — and because X801 has exactly flat magnitude, cascading them cannot double-correct anything. See [R7](#).

(Could a full complex inversion do both at once? Yes, and it would also try to invert the room’s excess phase — position-specific, acausal, pre-ringing. Separating loudspeaker excess phase from room excess phase needs frequency-dependent windowing of the kind DRC-FIR, Acourate and Audiolense do. That is a different tool, not a REW setting.)

4. The one law you cannot escape

$$\Delta f \cdot \Delta t \approx 1$$

A feature Δf wide in frequency takes about $1/\Delta f$ seconds to unfold in time. It is not a rule of thumb, it is the Fourier transform. It appears in this procedure in four disguises, and they are all the same statement:

you choose	you have thereby chosen
frequency resolution of the correction curve	the length of the filter’s ringing
the length of the IR window	the frequency resolution
the number of FDW cycles	the highest Q the correction may contain
smoothing bandwidth	the highest Q the correction may contain

Made concrete at 50 Hz:

resolution of the correction curve	width at 50 Hz	shortest possible ringing
1 FFT bin (no smoothing at all)	0.37 Hz	2731 ms
1/48 octave — REW’s “Var” below 100 Hz	0.72 Hz	1385 ms
1/12 octave	2.89 Hz	346 ms
1/6 octave	5.78 Hz	173 ms
1/3 octave	11.58 Hz	86 ms

Read the right-hand column as “*how long the woofer keeps moving after the music stops*”. **You do not get to choose the correction’s resolution and its ringing separately. They are one number.**

This is why “just correct everything accurately” is not available. The finer the correction, the longer the filter rings — and beyond some point you are trading an inaudible frequency-response error for an audible time-domain one.

5. The failure this guide exists to prevent

Run the inversion with defaults and this is what you get. It happened here, in September 2025, and the filters were in service for months.

The left channel’s measurement contains a **38.3 dB null at 28.93 Hz** — a razor-thin canyon, 2 FFT bins wide. It is not a room mode: the right channel varies only 6.4 dB over the same span, so it is one loudspeaker’s path to one microphone point. Nothing was regularised, so the inversion did its job faithfully and built the exact inverse: a **Q = 40** resonator.

what the deployed filter did	
narrowest feature below 200 Hz	0.8 FFT bins
group delay excursion, 20–200 Hz	80 ms (a correction filter should show a few ms)
28.7 Hz gated note, time to fall 40 dB	1348 ms

Audible as *the music stops and the woofers keep moving*. And the tell that it was never real: a Q of 40 at 29 Hz implies a 3.07 second decay. In 58 m³ that is physically impossible. **The filter was correcting something that does not exist.**

Every non-default setting in Part III is aimed at this one failure.

Part II — Decisions to take before you measure

6. One microphone position, or several?

Multiple positions is the only change on the list that attacks the cause rather than the symptom. Everything else in this guide suppresses narrow features on the grounds that they are *probably* not real. Spatial averaging finds out.

Why it works, and why it works especially well here

A genuine room mode is a property of the volume: a 50 Hz mode is at 50 Hz everywhere in the room, and only its *amplitude* varies with position. An interference null is a property of a path: two arrivals cancelling at one point, and it moves or vanishes when you move.

Average several positions and modes survive while interference washes out. That is a **discrimination no amount of smoothing can perform** — smoothing blurs everything equally, because a single measurement contains no information about which features are which.

And deep nulls are the most position-sensitive feature there is. A 38 dB null requires two arrivals to cancel to within a fraction of a decibel; move 20 cm and that balance is gone. So even a modest spread attacks exactly the feature that causes the damage.

The caveats, honestly

1. **Use RMS average for different positions. Never use Vector average across positions.** REW’s help: > “**Vector average** ... averages the currently selected traces taking into > account both magnitude and phase. It ... is most appropriate for multiple > measurements taken from the same position, or measurements which have been > time and level aligned.”

Vector-averaging *different* positions lets the position-dependent phase cancel and **manufactures new nulls** — the exact disease you are trying to cure. The **RMS average** is a power average, so a null at one position is filled by the others; that is the behaviour you want. **RMS + phase avg.** is not a better spatial average. Its magnitude is the one you want, but the phase it attaches is a vector average, which across positions means nothing — and REW warns it breaks the magnitude/phase relationship badly enough to need larger left windows (3d quotes it). The inversion generates minimum phase anyway, so use the plain **RMS average** action.

2. **Level differences between positions are weighting, not response.** REW suggests removing them with **Align SPL...** before averaging: *> “If the measurements were made at different positions (spatial averaging) > it is usually best to first use the **Align SPL...** feature to remove > overall level differences due to different source distances.”*

REW says no more than that, and how much it matters depends entirely on how far apart the positions are — [3b](#) measures it for a one-seat cluster and derives the rule.

3. **Spacing buys you less at low frequency than you would like.** Two points decorrelate when they are a useful fraction of a wavelength apart:

	30 Hz	50 Hz	100 Hz	200 Hz
λ	11.4 m	6.9 m	3.4 m	1.7 m
$\lambda/4$	2.9 m	1.7 m	0.86 m	0.43 m

A ± 30 cm cluster around one seat decorrelates well above ~ 250 Hz, partly from 100–250 Hz, and hardly at all below 60 Hz. **Spatial averaging does not solve the bass by itself.** It still helps below 60 Hz — via the null-depth sensitivity above — but do not expect the 30 Hz region to be cleaned up by moving the mic 30 cm.

4. **It is not free.** 5 positions $\times$ 2 channels = 10 sweeps, and one contaminated position — mic bumped, chair moved, room state changed halfway — poisons the average silently.

How far apart? Measure it, do not guess

A spread wider than the seat optimises a place nobody’s head goes, and a spread narrower than the seat does nothing at all. Measured in this room: how much the mono response actually changes when the microphone moves, rms over 1/6-octave bands, against the centre point.

microphone offset	20–80 Hz	80–225 Hz
8 cm sideways	0.40–0.44 dB	0.52–0.56 dB
10 cm forward	0.48 dB	0.64 dB
33 cm back	2.02 dB	1.70 dB

Below ~ 10 cm there is almost nothing to average. The $\lambda/4$ table above says why: at 50 Hz a decorrelating distance is 1.7 m, and 8 cm is 5% of it. Useful decorrelation starts somewhere around 20–30 cm, which is also where the listener’s head actually goes.

Across the three offsets available here the difference grows roughly **linearly with distance** — about **0.09 dB rms per centimetre** at the FDW’s own resolution:

offset	8 cm	10 cm	20 cm	33 cm
difference from centre	~ 1.1 dB	~ 0.9 dB	~ 1.8 dB (interpolated)	~ 3.0 dB

So ± 20 cm carries about twice the information of ± 8 cm, and ± 8 cm is close to nothing. That, and the head envelope above, is why the recommendation is 20 cm and not 8.

(One caveat this dataset cannot settle: its only large offset is also its only large front-to-back offset, and front-to-back is the axis that changes the microphone’s distance to the boundaries and so moves in-band interference. Distance and axis are confounded here. Measuring ± 20 cm on both axes separates them.)

20 cm is smaller than most guides recommend, on purpose

Mainstream practice assumes you are correcting a **listening area**. Dirac Live’s arrangements run to 9, 13 and 17 positions; the moving-microphone literature and most REW community guides suggest a cluster of roughly 0.5 m radius. Spread that wide and you decorrelate far more than the table above, which is exactly the point when several people are listening.

The cost is stated plainly in those same sources: a wide spread over-smooths, and under-corrects the seat you actually sit in. **This is a one-seat system**, so the cluster is sized to one head. If you want

the robustness of the wider spread, take it — but understand you are buying insurance against a head position nobody occupies, and paying for it in accuracy at the one that is.

The physics does not change either way, and nor does anything in Part III. Only the number of sweeps and the spacing move.

A useful independent check: the moving-microphone method. Play pink noise, set the RTA to a long-term (“Forever”) average, and walk the microphone slowly through the listening volume. It converges to the same magnitude as an RMS average of many discrete captures in that volume, faster and more repeatably, which makes it a good way to confirm that your five-position average is not an accident of where the five points happened to land.

It cannot replace the sweeps here. MMM yields **magnitude only** — no impulse response — so it can carry neither the FDW of step 2 nor the position-wise complex sum of 3c. Use it to check the answer, not to produce it.

Size the cluster from how the listener actually sits

On this sofa the head moves roughly **20 cm front-to-back** — slouching, and one or two cushions behind the back push you that far forward — and rather less sideways, though 20 cm each way still lands inside the seat.

That is not a small effect. The 2026-04-28 sweeps were taken about 20 cm behind the 2026-08-10 ones — both perfectly normal seating — and below 200 Hz the two differ by **3 to 6 dB**, with the room’s upper-bass dominance moving from 4.1 dB to 2.8 dB. **A normal shift in how you sit changes the bass more than any panel move measured in this room, and more than the choice of house curve.** That is the argument for averaging: not statistical hygiene, but the fact that a single-point filter is tuned to a posture.

Recommended

Five positions, all at ear height: the listening position and four points **20 cm** away from it — forward, back, left, right.

positions	5 — C, F20, B20, L20, R20. 20 cm is the smallest offset that measurably decorrelates in this room, and it is inside the seat
per channel	yes, L and R separately at every position; never average L with R at this stage
sweeps	10, plus one optional simultaneous L+R at C
combine with	RMS average across positions. Level alignment is optional at this cluster size — 3b has the measurement
then	build spatial L, spatial R and the spatial mono-sum divisor as step 3 prescribes; the sum must be formed position by position, before spatial phase is discarded

Keep everything except the microphone fixed for the whole set: speaker positions, gain, room state, mic height and orientation. Record the actual offsets, and check them against the front-wall distance REW writes into each measurement note — a mis-stepped position is invisible in the response and obvious in the distance.

If you stay single-point

Survivable, and this is the configuration the numbers in this guide come from. **Frequency-domain regularisation is the single-point substitute for spatial averaging** — both suppress narrow position-specific structure; the FDW just does it without moving the microphone. Accept these:

1. **The FDW is mandatory, not advisory.** With a spatial average you could argue for 20 cycles; with one point you cannot.
2. **Never boost a narrow null.** A deep narrow null is destructive interference between two arrivals — it is **not minimum phase**, so a minimum-phase inverse is the wrong *shape*, not merely the wrong size. The 0 dB max-gain setting enforces this; do not raise it.
3. **Keep the L/R average as the target** (step 5). It is already a partial average and it stops each channel chasing its own private nulls.
4. **Trust nothing below ~35 Hz.** The N801s are –6 dB around there anyway; correction buys excursion, not output.

5. **Above Schroeder (~166 Hz) single-point is fine** — the response there is direct sound and early reflections, which barely move over a seat.

The honest cost: a single-point correction is about as good as a multi-point one *at the microphone*, and noticeably worse one seat over. **It will not ring**, which is the thing you are fixing.

7. With several positions, do you still need the FDW?

Yes. Less of it, but you cannot drop it. This is the question people get wrong, so it is worth being precise about *why*.

Spatial averaging and time windowing answer **different questions**:

	spatial averaging	FDW
the question it answers	<i>is this feature real, or is it this seat?</i>	<i>how long may the correction be allowed to ring?</i>
removes	features that differ between positions	energy arriving late
what it does to a genuine high-Q room mode	nothing — it is at the same frequency everywhere	caps its Q at N

That second row is the whole argument. **A room mode survives spatial averaging with its Q intact** — that is what makes it a mode rather than an artefact. If the average still shows a $Q = 25$ peak at 45 Hz, it is real; and inverting it still builds a filter that rings

$$T_{60} = 2.2 \cdot Q / f_0 = 2.2 \times 25 / 45 = 1.2 \text{ seconds.}$$

Averaging established that the feature is genuine. It said nothing about whether correcting it is a good idea. **Only the FDW bounds the filter**, and §4's conservation law makes that inescapable: no number of microphone positions can make a 1 Hz-wide correction shorter than a second.

What multi-position buys you is permission to relax the FDW, because it is no longer doing two jobs at once:

measurement	FDW	reasoning
single point	12 cycles	must suppress interference <i>and</i> cap Q
5-position RMS average	15–20 cycles	interference already suppressed; only the Q cap is left
9+ positions, well spread	20 cycles	as above, with more confidence

Never off, and never above ~25 — beyond that you are no longer regularising anything and §5 becomes available again.

Measured: the two really are independent

The intuition that windowing and averaging are two denoisers doing the same job, so that stacking them over-cleans the data, is worth testing rather than assuming. Two measurements settle it.

The FDW does not remove what averaging removes. How far each position sits from the centre point, mono sum, 20–225 Hz rms:

offset	raw	FDW 12	FDW 20
8 cm lateral	0.73 dB	1.17	1.09
8 cm lateral	1.01	1.07	1.20
10 cm forward	0.86	0.86	0.84
33 cm back	2.71	2.96	2.94

Windowing leaves the position differences unchanged or slightly **larger** — §8's own mechanism running in your favour, since gating late energy exposes the early interference that is precisely what moves when the microphone moves. Whatever spatial averaging is removing, the FDW has not already removed it.

And the FDW does not shift the tonal target. The target is built from the same windowed traces as the divisor, so the window largely cancels in the ratio. Building the filter from windowed traces and from raw traces, 20–225 Hz:

+0.01 dB mean 0.56 dB rms 1.95 dB max

The broad tonal balance is identical to a hundredth of a decibel; only narrow structure changes, which is the entire purpose. The concern that windowing makes you “*EQ toward a target the room does not deliver*” is real for a workflow that windows the divisor and not the target — this one windows both.

The direct-sound effect people have in mind is real, but it lives above the correction band. FDW 12 minus raw, per octave:

20–32	32–63	63–125	125–250	250–500	500–1k	1–2k Hz
–0.6	–1.1	–2.1	–1.3	–2.2	–3.1 / –5.4	–2.7

Inside 20–225 Hz the window costs 1.4–1.6 dB overall and room gain survives. Above 250 Hz it starts genuinely stripping the reflected field — which is one more reason the correction band stops at 225 Hz.

8. How much regularisation — choosing the FDW cycles

The **frequency-dependent window** is a Gaussian time window whose width varies inversely with frequency. Its width is set in **cycles**, and the conversion is exact and worth memorising:

A window of **N cycles** is N/f seconds wide at frequency f , so its bandwidth is f/N and its **fractional bandwidth is exactly $1/N$** . Therefore **N cycles** $\leftrightarrow$ **a cap of $Q = N$** .

FDW setting	fractional bandwidth	$\approx$ octave fraction
15 cycles (REW default)	1/15	$\sim 1/10$ octave — mild
12 cycles	1/12	$\sim 1/8$ octave — the starting point; see the N = 8 callout below
~ 9 cycles	1/9	1/6 octave
~ 4.3 cycles	1/4.3	1/3 octave

(a) the FDW is not one window but a *family*, one per frequency, of width N/f — at 12 cycles that is 480 ms at 25 Hz, 240 ms at 50 Hz, 60 ms at 200 Hz. (b) the same windows as frequency kernels: **equally wide on a log axis**, which is what constant fractional bandwidth means and why the FDW behaves like fractional-octave smoothing. (c) the real measurement through it at 4, 12 and 30 cycles.

The three constraints that pin N down

1. The floor — do not truncate real modal decay. N must exceed the Q of the modes you intend to keep. From the room’s measured decay (R6):

at	credible T60	implied Q	cycles needed
50 Hz	~ 400 ms	9.1	9
80 Hz	~ 350 ms	12.7	13
120 Hz	~ 300 ms	16.4	16

$\rightarrow N \gtrsim 12$.

2. The ceiling — refuse position-specific interference. The features that wrecked the September filters had $Q = 40$ and $Q = 63$. N must be far below them.

$\rightarrow N \ll 40$.

3. The measurement must actually contain the window. At 25 Hz, $N = 30$ asks for 1200 ms of clean decay above the noise floor. It is rarely there.

N = 12 sits inside all three, and the floor and ceiling are a factor of 3 apart — this is not a knife-edge setting.




Figure 1: The FDW in time, in frequency, and applied to the measurement

The floor is about modes. It does not apply to a cancellation — use $N = 8$

Constraint 1 says $N \gtrsim 12$ because N must exceed the **Q of the modes you intend to keep**. A speaker-boundary null has no Q to keep: it is two paths cancelling, not energy being stored and released. The floor simply does not bind on it, and where the sharpest feature in the band is a cancellation rather than a mode, **$N = 8$ is the better setting**.

Built both ways from the same 2026-08-10 measurements, everything else identical:

	FDW 12	FDW 8
74 Hz null, depth below its 60–95 Hz shoulders	−27.8 dB	−15.6 dB
narrowest feature, L / R	31.6 ✓ / 25.9 × bins	49.6 ✓ / 49.6 ✓
group delay, L / R	+24.1 / +24.5 ms	+17.2 / +17.4 ms
gated tones failing, L / R	2 / 5	1 / 0

The raw measurement puts that null 18.8 dB below its shoulders. **A 12-cycle window made it 27.8 dB — nine decibels deeper than it really is**, by stripping the late energy that fills it (the warning in “Two things the FDW is not”). The filter then built correspondingly steep shoulders, and those shoulders were the 63 and 79 Hz tails.

What it cost, measured as the change to the correction curve over 20–225 Hz: **0.58 dB rms on L, 0.77 on R**, and spent where you would want it — −1.4 dB at 74 Hz, +1.2 at 80, −1.3 at 100, and **0.0 to 0.3 dB below 63 Hz**. The §11 bass work is untouched and the delivered mono sum moved 0.11 dB or less in every band bar 160–225.

So the rule is not “12 is wrong” but: **N is chosen against the sharpest feature you intend to correct**. Identify what that feature is first. If it is a mode, constraint 1 applies and $N \gtrsim 12$. If it is a boundary cancellation — which this room’s is, at 74 Hz, predicted by geometry at 75.3 Hz — the floor is not in force and a shorter window is both more honest about the measurement and kinder to the filter.

What 12 cycles does to the feature that caused the trouble

The 38.3 dB razor null at 28.93 Hz:

	level at 28.93 Hz	depth below the 27.10 Hz peak
raw, no FDW	33.7 dB	38.3 dB
FDW 30 cycles	50.7	20.4
FDW 12 cycles	58.5	11.7
FDW 8 cycles	62.8	6.9
FDW 4 cycles	68.3	1.3

12 cycles turns a 38 dB canyon into a 12 dB dip — no longer something an inversion will build a Q-40 resonator to fill — while the surrounding modal structure survives intact.

Two things the FDW is not

It is not magnitude smoothing. It is a time gate, so it changes the **complex** response. Where late energy had been *filling* a direct-sound null, tightening the window can make that null **deeper**, not shallower — visible on the 4-cycle trace near 70 Hz in panel (c). The 28.93 Hz null shallows because it is late-arriving interference; a genuine early-arrival null would not.

It is not destructive. The window is a *derivation setting* over a retained impulse response: untick it, press Apply Windows, and the original response returns exactly. What is **not** reversible is anything you already derived from it — see step 2.

9. Which smoothing, and why the smoothing menu will not help you

Set the smoothing menu to whatever you find easiest to read. It will not reach the arithmetic. This is the single most important practical fact in the procedure, and it is documented:

*“The result of arithmetic on measurements that have compatible impulse responses is smoothed using the measurement A smoothing, **unsmoothed data is used during the calculations.**”* — REW help, All SPL Graph → Trace Arithmetic

That sentence reads as self-contradictory and is not. It joins two clauses about different objects: the **output** trace inherits A’s smoothing *setting* for display; the **input** to the math is the raw impulse response. Compute on raw data, then hand the result A’s display setting.

So the truthful summary is:

REW operation	uses smoothed data?
text / log export	yes — “the smoothing for the exported data”
EQ target match (Auto EQ)	yes — “apply ... before running the target match”
trace arithmetic on IR-compatible measurements	NO — unsmoothed, always
trace arithmetic on non-IR traces (e.g. 96 PPO imports)	yes, whatever they carried

Applying 1/6 octave to LX-MP and pressing divide **does nothing at all**. REW reaches past it to the impulse response.

What does reach the arithmetic is the window, because the window *defines* that impulse response. REW says so of the FDW specifically:

*“Any frequency-dependent settings are excluded, applying an FDW to the result would amount to applying the window twice, **as it is already applied to the data used to produce the result.**”*

That is why this procedure is built around the FDW and not around the smoothing menu.

If you ever do need smoothing to stick

Round-trip it: apply the smoothing, File → Export → Measurement as text **with that smoothing selected in the export dialog**, then re-import. The exported numbers carry the smoothing, and on re-import they *are* the data. This is redundant if the FDW is set, and is listed as optional in step 4a.

“Smooth the spatial average instead of windowing the captures”

A reasonable-sounding alternative: average the raw captures, then apply Var or 1/6-octave smoothing to L-SP and R-SP to tame the residual ripple. It does not work here, for two reasons that stack.

Smoothing does not reach the division. That is this section’s whole subject: REW computes trace arithmetic on the impulse response and hands the result A’s smoothing only for display. You would be looking at a smooth trace and dividing by a rough one.

And it cannot cap the filter’s Q anyway. Smoothing a magnitude average is a frequency-domain operation on a trace that has already lost its time information; the FDW is a time gate applied while the time information still exists. §7 measures the difference, and step 2 measures what happens if you swap the order.

The advice is sound in the workflow it comes from — deriving an EQ *target* for REW’s own filter matcher, which carries filter-count, max-Q and max-boost guards of its own. A bare inversion has none of them, which is what §5 looks like.

And do not use REW’s recommended smoothing here. “Variable smoothing is recommended for responses that are to be equalised” is advice for the **target-match / parametric** path, which carries its own filter-count, max-Q and max-boost guards. A bare inversion has none of them. Worse, Variable is **1/48 octave below 100 Hz**, which at these frequencies is not smoothing at all:

freq	1/48 oct width	in FFT bins (0.3662 Hz)
20 Hz	0.29 Hz	0.79
25 Hz	0.36 Hz	0.99
28.9 Hz	0.42 Hz	1.14
51.3 Hz	0.74 Hz	2.02

freq	1/48 oct width	in FFT bins (0.3662 Hz)
100 Hz	1.44 Hz	3.94

A kernel narrower than one bin is an identity operation. Below ~50 Hz, selecting 1/48 octave is identical to selecting None.

For *looking* at the response while you work, plain **1/6 octave** is the right choice — it shows you roughly what the FDW is doing to the data. Avoid **Psychoacoustic** for anything but viewing: it uses a cubic mean, a deliberately peak-biased estimator worth up to **+1.03 dB** on this dataset, which becomes over-cutting if it ever reaches an inversion.

10. The correction band

20 – 225 Hz.

edge	why
20 Hz low	below it the N801 is rolling off and the room measurement is mostly noise; correction there buys woofer excursion, not output
225 Hz high	above Schroeder (166 Hz) the response is direct sound plus early reflections — a property of the loudspeaker and the seat, not of the room. It does not need a 128k-tap filter, and correcting it from one microphone point fits that point

The 225 Hz figure gives about half an octave of overlap above Schroeder, which is enough to let the correction die away gracefully rather than stop dead. REW blends the limits over **one octave**, so the correction is effectively fully out by ~320 Hz and fully in by ~160 Hz.

Why not push the top edge higher

The honest objection to 225 Hz is that the response above it is *not* featureless. On the five-position set (`../DRC-120.blue/120.blue.multi.pt.mdat`) two features sit just above the edge and are plainly visible in `LR.filtered`:

feature	in the spatial average LR-SP-MP
valley, 226 – 410 Hz	–5.2 dB at 312 Hz (re its own 200 Hz – 2 kHz mean)
bump, 500 – 726 Hz	+4.6 dB at 552 Hz

The first thing to check is whether they are single-point interference. They are not — or not entirely. Normalising each of the five positions to its own 200 Hz – 2 kHz mean, 1/6 octave:

band	mean over the cloud	σ across ± 20 cm
288 – 367 Hz	–2.4 ... –4.8 dB	2.2 – 3.6 dB
524 – 609 Hz	+3.2 ... +5.1 dB	1.8 – 2.1 dB
900 – 1200 Hz	–4.3 dB	3.5 – 4.4 dB

Above 900 Hz the spread overtakes the feature and it is comb, exactly as §6 predicts. At 300 – 600 Hz the feature still stands at two to three times the positional spread, and it survives the five-position average. So “it is only comb” is not the reason to stop at 225 Hz.

The reason is the clamp. Carry it through `min(Target – divisor, 0)`, in raw dB, with the target level REW itself calculated:

f	LR-SP-MP	Target	LR.RMS	cut delivered
303 – 345 Hz	63.4 – 64.3		71.6	0.00
470 Hz	67.8		71.5	0.00
536 Hz	72.8		71.5	–1.32
560 Hz	73.1		71.5	–1.63
585 Hz	72.5		71.5	–1.01
≥ 611 Hz	≤ 71.1		71.5	0.00

- **The valley is a dip, and the correction is cut-only.** It clamps to exactly zero across the whole 226 – 410 Hz band. There is nothing to decide here: §4 forbids filling it, and the arithmetic does not even try.
- **The bump yields –1.6 dB at 560 Hz**, over roughly 536 – 600 Hz. Not the 4.6 dB it looks like on a normalised plot — the target line sits about 3 dB above the response’s own 200 Hz – 2 kHz mean, so only the tip of the bump clears it.

Extending the per-channel band from 225 Hz to 800 Hz therefore buys **one 1.6 dB trim half an octave wide** ($Q \approx 2$, so it would not trouble R8’s sharpest-feature test — it simply is not worth a rebuild).

The trap. The bump *looks* like 4.6 dB, and the tempting fix is to raise the target level until the clamp bites harder. Do not. Raising it cuts broadband to chase a shape — the broad –3 to –6 dB roll above 600 Hz — that is the loudspeaker’s own power response and the seat, not the room. You would spend headroom to install a treble tilt you did not ask for.

What the features actually are. Not one comb: the peaks land at 216, 415, 575 and 705 Hz, spacings 199, 160 and 130 Hz — decreasing, so no single reflection delay fits. Gating the centre-position IR, the 297 – 362 Hz valley is absent at a 3 ms gate and only appears from 5 ms out; the ETC shows a –17 to –21 dB plateau from 3.5 to 7 ms, which is the side wall, the ceiling and the front wall arriving together. That is **H(excess)** — interference between direct sound and early reflections — and §2 says what a minimum-phase filter does to it: nothing. Moving it at one point moves it somewhere else.

The lever for that band is acoustic, not arithmetic. Absorption at the first reflection points acts on it at every seat, which no EQ can. On this geometry two front-wall screens delivered –2.0 dB over 355 – 710 Hz at roughly 5σ against head movement, and the intervention was judged under-dosed; see the dated measurement notes for that build.

One caveat left open. The valley bottoms at 303 – 345 Hz, which is the N801’s woofer-to-midrange crossover at 350 Hz. If part of it were the crossover sum it would be minimum phase and genuinely correctable — but it would need a *boost*, which cut-only forbids, so it changes nothing operationally. The near-field material in `../DRC-120.blue/on.axis.61cm.txts/` cannot settle it: those exports start above ~500 Hz.

11. Below 80 Hz, correct the sum — not each channel

§5 warns against letting each channel chase its own nulls, and fixes it by giving both channels a **common target**. That is necessary and it is not sufficient. The divisor is still each channel’s **own** response, and below about 80 Hz that is enough to do real damage.

The failure this section exists to prevent

Measured on the 2026-08-10 Rscreen pair, at 50 Hz:

	L alone	R alone	RMS average	vector average = normalized mono sum
level	75.92	71.83	73.19	68.45 dB

The two speakers **cancel each other** at the listening position across 45–56 Hz — the sum is 4.7 dB *below* the average of the parts. The RMS average cannot show this: it is a power average, it has no phase, and filling nulls is exactly what it is for.

So the divisor said “73.2 dB, target 71.4, cut it”, and the build cut **L by 4.20 dB and R by 0.00 dB**. But L was the *louder of two nearly anti-phase contributors*. Making them more equal made the cancellation more complete:

	before	after the per-channel build
mono sum at 50 Hz	68.45	65.96 dB
40–62.5 Hz, re midrange	–0.90	–3.49 dB

The region that was already the weakest in the response became the deepest hole in it, and the filter did that — correctly, by its own arithmetic.

The one-line reason

If both channels receive the **same** correction C , then **new sum = sum + C exactly**: the ratio between the two contributions never changes, so the depth of any cancellation between them is untouched. Only a **differential** correction can deepen a mono null.

Per-channel division against a common target is differential by construction. Above the transition it is what you want — it is how each speaker's own response gets flattened. Below it, it is unforced error.

The fix is the divisor, not the target

The target does not change. Keep the EQ-window target built from the plain L/R RMS **average** exactly as §5 prescribes — a null in one path filled by the other, one common shape for both channels. Change only **what you divide by**, and only below 80 Hz:

Hz	target	cut asked vs RMS average	cut asked vs mono sum
20	69.79	4.98	5.05
25	72.21	3.99	4.27
31.5	72.56	1.34	0.62
40	72.17	0.00	0.00
45	71.81	0.00	0.00
50	71.36	1.65	0.00
56	70.89	2.78	0.00
63	70.20	2.73	2.40
71	69.53	4.27	3.55
80	68.84	7.19	5.17

Against the sum the filter asks for **nothing** at 45–56 Hz — not because anything was special-cased, but because the sum is already below target there and **the cut-only clamp engages by itself**. The deep-bass trim at 20–31.5 Hz survives intact. This is the whole technique: give the division a divisor that knows about the cancellation, and the existing guard does the rest.




Figure 2: Why the divisor must be the sum below 80 Hz

Where 80 Hz comes from

Two independent arguments land on the same number.

Coherence. Below ~80 Hz the two speakers arrive at the seat correlated, so they sum as pressure and the sum is the physical quantity. Above it they decorrelate and the vector average stops meaning anything — measured here, RMS average minus vector average:

20	31.5	40	50	56	63	125	200 Hz
−0.07	+0.71	+1.23	+4.61	+4.55	+0.33	+0.16	−2.75

They agree within a decibel everywhere except the cancellation — and by 200 Hz the vector average has fallen 2.75 dB below, which is decorrelation, not information. **Use the sum only where it means something.**

How differential the correction is anyway. Channel minus average, rms:

	20-40	40-63	63-80	80-160	160-225 Hz
L	1.22	2.68	1.23	1.16	2.87
R	2.09	3.09	2.00	1.36	1.87

The differential is worst exactly in the cancellation band and modest elsewhere, so a crossover at 80 Hz removes the harm and forfeits almost nothing. This is also the standard bass-management crossover (Dolby, THX, ITU), chosen for the same reason: localisation cues are weak below it.

Building it: two divisions, spliced at 80 Hz

$$FL = \left[\begin{array}{c} \text{Target} \div LR \\ \text{limited 20-80 Hz} \end{array} \right] \times \left[\begin{array}{c} \text{Target} \div LX \\ \text{limited 80-225 Hz} \end{array} \right]$$

Both factors are the same kind of object — **the target divided by a measurement** — and both are ordinary cut-only divisions. Outside its own limits each reverts to unity, so only one is ever doing work:

	below 80 Hz	above 80 Hz
Target ÷ LR , limited 20-80	active	unity
Target ÷ LX , limited 80-225	unity	active
product	Target ÷ LR — common	Target ÷ LX — per channel

The splice is seamless because REW’s two band-limit ramps are **complementary**: the roll-off at the upper limit of the first is exactly one minus the rise at the lower limit of the second, each a raised cosine over one octave. So in decibels the product is

$$FL(\text{dB}) = w \cdot (\text{Target} - LR) + (1 - w) \cdot (\text{Target} - LX), \quad w: 1 \rightarrow 0 \text{ over } 57\text{--}113 \text{ Hz}$$

— a clean crossfade from the common correction to the per-channel one. Nothing steps, and at no frequency is any correction applied twice.

Both factors stay cut-only — nothing has to be allowed to boost

This is the property that makes the construction safe without a post-hoc check. **Target** ÷ **LX** is a target-over-measurement division like any other, so **select Max gain and set 0.0 dB on both factors**, and cut-only is then structural rather than something you verify afterwards.

The equivalent form, and why it is the worse way to build it

The same filter can be written as a common correction times a *differential* one, $(\text{Target} \div LR) \times (LR \div LX)$, with the first factor spanning the full 20–225 Hz and the second limited to 80–225. **LR** then cancels algebraically above 80 Hz and the result is the same filter — measured on the 2026-08-11 data, the two agree to **0.088 dB rms and 0.425 dB maximum** over 20–225 Hz.

Prefer the spliced form anyway. **LR** ÷ **LX** is a ratio of two *measurements*, so it genuinely needs to boost wherever a channel sits below the sum, and clamping it breaks the algebra:

	$(T \div LR) \times (LR \div LX)$	$(T \div LR) \times (T \div LX)$
boost the second factor wants	+6.59 dB	+0.22 dB
Max gain value it must be given	+6.0 dB	0.0 dB
maximum of the product, 20–225 Hz	+0.20 dB	+0.00 dB
cut-only is	verified after the fact	structural

The differential form ends up 0.20 dB *above* unity — harmless in itself, but it means the cut-only guarantee is no longer something the settings enforce. The spliced form never leaves it.

What it delivers

Predicted mono sum, dB relative to the untouched midrange:

band	before	per-channel build	common < 80 Hz	recovered
20–40	+8.06	+6.15	+5.82	−0.33
40–62.5	+2.68	+0.32	+2.63	+2.31
62.5–100	+7.85	+2.22	+2.89	+0.67
100–160	+10.25	+1.85	+1.85	0.00
160–225	+6.96	+4.38	+4.38	0.00

The 40–62.5 Hz loss comes back essentially in full, and the correction that was doing real work — the 100–160 Hz lump, 8.4 dB of it — is untouched.

This is not a local invention

Correcting the summed low-frequency field rather than each source is the mainstream approach: **bass management** in every multichannel standard sums below 80 Hz; **Dirac Live Bass Control** jointly optimises all bass sources rather than each independently; **Multi-Sub Optimizer** optimises the summed response at the seats for precisely the reason above; **Trinnov**, **Audyssey Sub EQ HT** and **Anthem ARC Genesis** all apply shared correction below the crossover; **Harman’s Sound Field Management** optimises the combined field. The only unusual thing here is that the shared bass sources are the two main speakers rather than a subwoofer. The physics is identical.

Building L+R when only separate L and R sweeps exist

REW’s **Vector average** of a matching L/R pair reconstructs $(L + R) / 2$. Use it at each microphone position while coherent phase still exists, **before** spatial RMS averaging. Do not vector-average spatial L-SP and R-SP: those traces no longer contain a meaningful position phase.

The distinction between beta’s two similarly named actions matters:

action	result	use here?
Vector average	$(L + R) / 2$	yes — same level scale as one channel and the target
Vector sum	$L + R$	no — 6.0206 dB higher and needs explicit normalization

A simultaneous physical L+R sweep is also $L + R$, so subtract **6.0206 dB** before substituting it for a vector average. Use it un-normalised and you ask for roughly 6 dB of spurious common cut across the whole 20–80 Hz band — and **Max gain 0.0 dB** will not stop you, because it clamps boost, not cut.

Two independent confirmations of the figure. **The 185 cm set**, which contains a genuine measured L+R sweep beside its two channels: the complex sum of L0 and R0 reproduces the measured LR to **0.502 dB rms** over 20–250 Hz (median +0.010 dB), and that measured sweep sits **+6.031 dB** above the vector average — a different room, REW version and signal chain, landing on the same number.

And the 2026-08-17 centre set. Measured L+R sits a median **6.009 dB** above the calculated L/R vector average over 20–225 Hz, versus the theoretical 6.0206 dB. After normalization:

comparison, measured versus calculated	rms	median	maximum
80–225 Hz, native 0.366 Hz bins	0.224 dB	−0.006 dB	0.941 dB
20–225 Hz, 1/12 octave	0.504 dB	−0.007 dB	5.35 dB
20–225 Hz, 1/6 octave	0.276 dB	+0.003 dB	1.95 dB

The 1/12-octave maximum is the 49 Hz cancellation. At 49.072 Hz the separately measured pair predicts 75.06 dB while simultaneous $L+R/2$ measures 57.45 dB. This is not a broadband scale failure: a tiny sequential phase or response change moves an almost perfect cancellation between bins. Use simultaneous L+R at centre when available, and calculated vector averages elsewhere.

Spatial averaging makes the residual safe: in the five-position calculation, replacing calculated centre with measured centre changes the common response below 80 Hz by 0.174 dB rms and at most 0.97 dB at 1/12

octave; at 1/6 octave the change is 0.104 dB rms and at most 0.41 dB. The substitution is therefore sound for a spatial divisor, but it is not literally perfect at native resolution.

12. What a magnitude filter cannot do — and the one cheap way round it

Everything above corrects **magnitude**. There is one low-frequency failure it cannot address at all, and one inexpensive tool that can. This section is the worked case, from the 185 cm configuration in this repository.

The failure

Speakers 1.85 m from the front wall, microphone at 4.54 m, measured L+R:

deepest point of the mono sum, 20–300 Hz	– 17.9 dB re midrange, at 41 Hz
L alone, at its own worst	–8.3 dB
R alone, at its own worst	–10.0 dB

The sum is far worse than either channel, and it has two stacked causes:

	each channel already down	L and R cancel a further	total
36 Hz	–5.6 dB	–4.5 dB	–9.8
40 Hz	–6.0 dB	–7.3 dB	–13.1
43 Hz	–5.6 dB	–5.7 dB	–11.1
46 Hz	–3.7 dB	–3.5 dB	–7.0

The first column is **SBIR**: 1.85 m from the front wall gives an excess path of 3.70 m and a null at $343/(4 \times 1.85) = 46.4$ Hz, and it hits both channels identically because both sit at the same distance. The second is that L and R arrive at the seat **142° apart** at 40 Hz and partially cancel.

Neither is reachable by a cut-only filter, and boosting is not a serious option: +11.7 dB at 41 Hz is **15× electrical power** and **3.9× cone excursion**, spent on a null that moves with the listener’s head.

The half of it that is free

The SBIR component is geometry — nothing but moving the speakers fixes it, and that is what the 120 cm configuration is (it puts the null at 71.5 Hz).

But **the 7.3 dB of L/R cancellation is not a magnitude problem at all**. It is two sources at 142°, and rotating one of them costs no energy, no headroom and no excursion. A single second-order **allpass** does it. Measured — L0 through a 42.5 Hz, Q 2.5 allpass, summed with R0:

band	L + R	L_ap + R	change
20–32 Hz	+1.5	+1.6	+0.1
32–50 Hz	–11.8	–4.5	+7.3
50–63 Hz	+3.2	+3.2	–0.0
63–100 Hz	+6.0	+5.5	–0.5
100–225 Hz	+0.6	+0.5	–0.0

Deepest point –**17.9** → –**7.3 dB**; at 43 Hz the gain is +**11.7 dB**. Seven decibels recovered in the problem band, the rest of the range moved by less than half a decibel, and **not one decibel of boost applied anywhere**.

What it costs, and what it makes the filter

The measured allpass: magnitude flat to **0.26 dB** across 20–500 Hz, peak group delay **37.3 ms at 41.7 Hz**, falling to 1.9 ms by 80 Hz.

That 37 ms sits right at the low-frequency group-delay audibility threshold, and it is applied to one channel only. Below 80 Hz that is benign — the same reason bass management crosses there: localisation cues are weak, so an L/R group-delay mismatch has little to act on.

This makes the filter mixed phase — but the harmless kind

Any causal stable filter factors as *minimum phase* $\times$ *allpass*, and is minimum phase exactly when that allpass factor is trivial. Adding one makes it non-trivial, so §2’s “we restrict ourselves to minimum phase” no longer describes the chain.

It is still safe, because there are two different operations both called phase correction:

	what it does	causality	cost
add an allpass	re-times one source against another	causal	group delay only
invert the room’s excess phase	undoes an allpass the room applied	acausal	pre-ringing

The second is acausal by necessity — the inverse of a stable allpass has its poles outside the unit circle, so no causal stable inverse exists. That is the operation that needs windowing and that pre-rings.

This filter is the first kind. Its impulse response was checked: **16384 taps, peak at sample 0, exactly zero energy before the peak.** It can only delay. There is nothing in front of the impulse to pre-ring with.

What this is, in the wider landscape

Plain two-channel **Dirac Live** would not find this. It measures L and R separately and corrects each toward the target — which at 142° is exactly the operation §11 warns about, driving the cancellation *deeper*. The tools that find it are the ones that optimise the **summed** field across sources: **Multi-Sub Optimizer** (free, and the closest match to what was done here), **Dirac Live Bass Control / ART**, and **Trinnov** source optimisation.

REW can generate the filter natively: the EQ window offers an **All-Pass** filter type taking frequency and Q, and it exports into the impulse response like any other. No mixed-phase design software is required.

When to reach for it

Only when the mono sum is materially worse than either channel alone. Test it in one step: compare L+R against $\max(L, R)$. Coherent summation gives about +6 dB; anything at or below 0 dB means the channels are fighting, and that deficit — not the magnitude response — is what to fix first. Above ~80 Hz leave it alone: the phase relationship there varies too fast with position to be worth chasing, and localisation does depend on it.

Part III — The procedure

Eleven steps build and accept the filter (the diagram folds 10 and 11 into one box). Steps 1 to 3 turn ten sweeps into the three traces everything else divides by; steps 4 to 11 are the same whether you measured one position or five. A twelfth, [step 12](#), follows once step 11 accepts: preparing the export directory the open-media-drc deployment documentation and installer both expect.


Where the actions live. In V5.40 beta, select the traces you want in the **All SPL** legend, then right-click the graph: **Align SPL...**, **Cross corr align**, **Vector average**, **Vector sum**, **RMS average** and **RMS + phase avg.** are all on that menu. **Trace arithmetic**, **Generate minimum phase** and **SPL offset** are reached from the measurement’s own controls.

Worked numbers come from `120.blue.txts/` (single position) and `120.blue.screens.multimeas.mdat` (the multi-position set).

TL;DR — the eleven steps, briefly

For the impatient, or for a rebuild you’ve done before: what each step is *for*, in one to three lines, then the settings. Each step below links to its full version; read that if a line here doesn’t make sense yet.

1. **Measure.** An acoustic timing reference puts $t = 0$ at the impulse peak — required so the FDW (step 2) centres correctly and so the position-wise L/R sum (step 3c) means anything at all. Sweep each position,



Combine L with R only within a position (vector average); combine different positions only with RMS average.

Figure 3: The REW inversion chain

start ≈ 12 Hz (not 10 — woofer excursion), -12 dBFS ($-18\dots-20$ dBFS if starting below ~ 15 Hz). One reference speaker for every sweep of both channels.

2. **Window.** The FDW is the regularisation that keeps the inversion from turning a razor-thin, position-specific null into a Q-40 resonator (§5). It must be applied before any averaging — window-then-average and average-then-window are not the same operation, and there is no way back from getting the order wrong. Add FDW, **12 cycles**, **Apply to all**, **keep ref time**, on every original L and R capture. Never on X801, never re-applied to a derived trace. (*Beta: select traces in the All SPL legend, right-click the graph $\rightarrow$ RMS average / Vector average / Align SPL.... Stable: these live in separate Actions panels, not a right-click menu.*)
3. **Reduce to divisors.** The rest of the procedure needs exactly three traces: each channel's spatial average (for the per-channel filter above 80 Hz) and the mono sum's spatial average (for the common filter below 80 Hz, §11 — the two speakers sum coherently down there, so only the sum may be corrected). Five sub-steps, in order — arithmetic is a snapshot, so redoing an earlier one means redoing everything after it.
 - **3a**, optional. Repeat captures at one position are a repeatability check, not extra spatial samples: **Vector average** them down to one trace per channel per position *before* anything else sees them, or skip if you took one sweep each.
 - **3b. Do: nothing.** Do not run **Align SPL** — at a 20 cm cluster, the table in that sub-step shows doing nothing beats every alignment policy including **Align SPL** itself, so skip 3b's whole procedure and go straight to 3c. (**Align SPL** only becomes the right call once positions span a metre or more, which is not this project's cluster.)
 - **3c.** The one place inter-channel phase is used. **Do:** for each of the five positions, select its L and R capture and choose **Vector average**. Name the results LR C, LR F20, LR B20, LR L20, LR R20.
 - **3d.** Now discard spatial phase. **RMS average** each channel's five positions $\rightarrow$ L-SP/R-SP; **RMS average** the five LR traces $\rightarrow$ LR-SP. Never **Vector average** across positions — it manufactures new nulls. (*-SP = spatial average — it names the dimension being averaged over, the five mic positions, not what survives the average. That it's an RMS (magnitude-only) average is a separate fact, stated by "discard spatial phase" above: -SP flags "already averaged across positions," distinct from a plain L/R/LR, which is one position and — for LR specifically — still carries the phase 3c built it with.*)
 - **3e.** Import X801.wav (confirm it wasn't already loaded before step 2's **Apply to all**; offset -117 dB $\rightarrow$ **Add to data** so it reads 0 dB, not the $\approx +117$ dB REW's import gives it with no calibration reference to work from). Then $LX = L-SP \times X801$, $RX = R-SP \times X801$ — the system as it will actually play, crossover included, before inversion.
4. **Minimum phase, first time.** Converts a magnitude-only divisor into a causal impulse response, so dividing by it later doesn't ask REW to invent phase from nowhere. $LX/RX/LR-SP \rightarrow -MP$ copies. **Cal file effects included** — this is modelling the acoustic response as measured, and the mic calibration is part of what that means (contrast step 8, where it's excluded: a filter has no microphone). LF tail on, corner at the sweep start; slope 24 dB/oct (ported) or 12 if the corner sits within $\frac{1}{2}$ octave of the correction band's low edge — never 0 dB/oct, that's step 8's move, not this one. HF tail off.
5. **Build the target.** What you want the room to measure like — everything downstream is this divided by what it actually measures. A common target, not each channel's own smoothed response, so correction doesn't pull the stereo image. **RMS average** $LX/RX \rightarrow$ L-R **RMS average**. Open it in the **EQ window**; load a house curve there if you want one (no scoop if the room already runs full — `tools/housecurve/house-curve-harman-fuller.txt`). Press **Calculate** to set the target level from that response — REW anchors it to the speaker's midrange, which is the number to check it against, not the neighbouring bass. Move the target's LF cutoff to 5–10 Hz, below the correction band. Save the shape: this **generates a new measurement**, **Target L-R RMS average**, which is what step 7 divides by — magnitude only, no phase to speak of.
6. **Stop and verify.** A narrow dip in an -MP divisor becomes a narrow **boost** once you divide by it — the one place a defect is cheaper to catch than to build and then discover. Check each -MP copy against its source for narrow dips before dividing by it.
7. **Divide.** Target $\div$ measurement *is* the correction curve, by construction (§1). Splitting the divide at 80 Hz keeps the sum-only correction (§11) from ever reaching a per-channel trace, and the cut-only ceiling stops the inversion from boosting into a null it should instead be regularised away from (§8). **Target $\div$ LR-SP-MP limited 20–80 Hz** (or 25–80 if chasing the group-delay gate) $\rightarrow$ **Fcommon**. **Target $\div$ LX-MP / RX-MP limited 80–225 Hz $\rightarrow$ Fper_L/Fper_R**. Multiply: **Fcommon $\times$ Fper_L $\rightarrow$ FL** (and FR). **Max gain on, 0.0 dB** throughout — cut-only. (*Beta: Max gain, a dB ceiling. Stable: a Regularisation percentage instead — a different control, not a unit conversion of the same setting.*)
8. **Minimum phase, second time.** Turns the raw FL/FR correction curve into a realisable causal filter — the same reason as step 4, applied to the thing you're about to export rather than to a divisor. **FL/FR $\rightarrow$ LFilter/RFilter**. LF tail on; slope shallowest offered (beta goes to 0 dB/oct — flat, no imposed roll-off) unless you want a deliberate subsonic high-pass. **Cal file effects off**.

9. **Bake in the crossover.** X801 corrects the one thing a minimum-phase inversion is provably blind to — the crossover’s all-pass phase (§3) — and it must go last because it has nothing to interact with until the room-correction filter exists. $X801 \times LFilter \rightarrow FLX$ (and FRX), X801 as trace A, last of all.
10. **Export.** What BruteFIR actually loads; trimming only sets latency and does not change the exported response. Trim to set latency, export 48 kHz / 32-bit float WAV.
11. **Accept or reject.** The numeric form of “will the woofers keep moving after the music stops?” (§4’s $\Delta f \cdot \Delta t$ trade-off, made concrete) — the one check that catches what a magnitude plot cannot show at all (R8’s worked example). `tools/drc_acceptance.py` on both channels. Fails $\rightarrow$ back to step 2, not a post-process patch on the WAV.
12. **Prepare the export bundle.** Ten files, matching open-media-drc’s deployment/provenance requirements. The one to get right: L/R still carry step 2’s FDW at this point — `L.txt/R.txt/LR.txt` must instead come from **unwindowed duplicates** (untick FDW, re-export), because open-media-drc’s own automated check only catches REW’s **Smoothing** setting, not FDW, and would accept a windowed export silently. Set **Smoothing: None** on every export. Build `L.Filtered/R.Filtered/LR.Filtered` by multiplying those same unwindowed duplicates by the imported `FLX-trimmed/FRX-trimmed` (checking the import level, as step 9 did for X801).

Step 1 — Measure the five positions

Do:

Before the first sweep: level-match the two speakers at the source, with REW’s generator and an SPL meter, and then leave the volume path alone for the whole session. From here on the measured L-versus-R level difference is a property of the room and the seat, and 3b will refuse to touch it — so it had better not also contain an electrical imbalance you could have removed.

setting	value
sweep	512k log swept sine , one sweep
start frequency	$\approx$ 12 Hz — low enough that step 4’s LF-tail corner (which REW floors at <i>start + 1 Hz</i>) lands a clean octave below the 20 Hz band edge, without going so low it over-drives the woofer. See the caution below
level	−12 dBFS (drop to −18...−20 dBFS for any sweep starting below ~15 Hz)
timing	acoustic timing reference — required
reference speaker	the same one for every sweep of both channels. Measure L with <code>Output = L, Ref = L</code> ; measure R with <code>Output = R, Ref = L</code>
sample rate	48 kHz
channels	L and R separately , at every position
positions	C, and four points 20 cm away — F20 forward, B20 back, L20 left, R20 right, all at ear height
trace names	channel first, then position: L C, R C, L F20, R F20, ...
useful extra	one simultaneous L+R sweep at C, at the same per-channel generator level

That is ten sweeps, eleven with the optional sum. Work position by position — move the microphone once, take L then R — so that the two channels at a point always share the same microphone placement.

Between sweeps, change nothing but the microphone. Same speaker positions, same gain, same room state, same mic height and orientation. Write the actual offset into the measurement note, and check it against the front-wall distance: F20 and B20 must differ from C by 20 cm, L20 and R20 must not differ from C at all.

[!] How low to start the sweep — and why not 10 Hz

Step 4 needs the LF-tail corner an octave below the 20 Hz band edge, and REW will not put the corner lower than *sweep-start + 1 Hz*, so the sweep has to start around **12 Hz** to give the corner room. Lower buys nothing: there is no usable data below ~16 Hz — the room measurement there is noise and the speaker is rolling off — REW only needs the low start to place the corner. And a log sweep dwells **equal time per octave**, so 10–20 Hz gets as much excitation as 10–20 kHz; on a ported woofer, below port tuning the cone unloads and excursion climbs steeply for constant drive.

So: **start at 12–14 Hz, not 10**, drop the level to $-18\dots-20$ dBFS for those sweeps, and run one trial while watching the woofer — visible large excursion or port chuffing means raise the start or lower the level.

[!] One reference speaker for the whole set — this is what makes 3c possible

The acoustic timing reference does two jobs. It puts $t = 0$ at the impulse peak, which is what the FDW needs. And, *if every sweep of both channels references the same speaker*, it puts L and R on one common clock — which is the only reason the position-wise complex sum in 3c means anything.

Switch reference speakers between the L pass and the R pass and each channel gets its own arbitrary $t = 0$. The magnitude averages survive that; the mono sum does not, and it fails silently.

It is recorded in every export header, so it is checkable after the fact:

```
grep -o 'Acoustic reference played from [A-Za-z0-9 ]*' *.txt
```

Every line must name the same speaker. On the 2026-08-17 set all fifteen captures read **played from DAC8STEREO L**.

A second check, from the same headers. With the left speaker as the reference, `grep -o 'Delay [-0-9.]* ms' *.txt` should show two distinct populations: the **L captures sit at essentially zero**, because source and reference are the same speaker, while the **R captures carry the real left-to-right arrival difference at that seat** and swing with the lateral positions. On the 2026-08-17 set the R captures read -14 mm at centre, **+43 mm** 8 cm to the left and **-84 mm** 8 cm to the right — the right speaker getting farther as you move left, which is the geometry. If the R captures also sit at zero, the reference speaker was switched between passes and 3c has nothing to work with.

Microphone orientation: point it at the ceiling and load the 90° calibration file, so the capture is insensitive to exact aiming. In the 20–225 Hz correction band this changes almost nothing — the capsule is omnidirectional down there — but the 500 Hz–2 kHz band that 3b may use for level alignment is squarely in the region where aiming does matter, so pick one convention and hold it for the whole set.

Repeat sweeps are optional. Two or three captures at the same point test repeatability and reduce noise; they are *not* extra spatial samples, and [step 3a](#) collapses them to one trace before the spatial stage sees them.

Why the acoustic timing reference: it puts $t = 0$ at the impulse peak. The FDW is centred on the window reference time, and REW's help is explicit that “for best results this should be at the peak of the impulse”. Without a timing reference the reference time is an estimate and the window sits in the wrong place — quietly, with no error message.

Step 2 — Set the window on every original capture, before averaging

Do: open **IR Windows** on L. Change **one** field:

control	set to
Left window shape / width	leave (Rectangular, 500 ms)
Right window shape / width	leave (Rectangular, 1000 ms)
Ref Time	leave — it is the IR peak
Add FDW	✓ ticked, 12 cycles

Then press **Apply to all, keep ref time**. Apply it to every original L and R capture and to any measured L+R capture. Never apply it to X801, and do not put it back on a derived average.

[!] X801 must not be in the measurement list yet

Apply to all means all — every measurement currently loaded, with no way to exclude one from the selection. X801 is not imported until 3e, several sub-steps later, precisely so this button cannot reach it. If you are re-running this step on a session where X801 was imported early — for reference, or left

over from a previous build — remove it from the measurement list (or at least deselect it) before pressing **Apply** to **all**, then re-import it fresh at 3e. Windowing X801 is not a small error: see “Never window X801” below step 3e for what an FDW does to an all-pass designed to be unwindowed.

[!] The order is load-bearing: window first, average second

The tempting alternative is to average the raw captures and apply the FDW to L-SP and R-SP afterwards. It is not available, and it would be wrong if it were.

It is not available. RMS average discards phase — REW: “*Phase is not taken into account, measurements are treated as incoherent... the result has the magnitude data from the source measurement and no phase data.*” No phase means no impulse response, and an FDW is a window on an impulse response. REW also clears the flag on derived traces: “*applying an FDW to the result would amount to applying the window twice.*”

And it would be wrong. The FDW discriminates in *time*. Each position has its own impulse response with its own arrival times, because the microphone moved. Window each capture and every reflection is gated where it actually lands; average first and those arrivals have been merged into a magnitude ripple with no time coordinate left, so a window applied afterwards gates an artefact. The two orders are not the same operation — measured on the five L captures at 12 cycles:

window-then-average vs average-then-window	
20–40 Hz	0.26 dB max
40–80 Hz	4.01 dB max
80–150 Hz	0.68 dB max
150–225 Hz	13.96 dB max
20–225 Hz overall	2.05 dB rms

Window each of the ten original captures. Only then form any average. There is no way back from getting this the wrong way round except to redo the chain.

Why only the FDW: at 12 cycles the FDW is narrower than the 1 s right window everywhere above 12 Hz, so it governs the entire working band on its own. The 1 s rectangular truncation is harmless because the decay is already in the noise by then. Shortening the rectangular right window instead would manufacture ripple (R2); this route avoids the question entirely.

Why keep ref time and not Apply windows to all: each measurement has its own reference at its own IR peak. Plain “apply to all” overwrites those with *this* measurement’s Ref Time and misaligns the others. “Keep ref time” copies the shapes and widths and leaves each measurement’s own peak alignment intact.

L and R showing different Ref Times is correct, not a fault

Every measurement carries its own Ref Time because every sweep began at a different instant in its own capture. What matters is whether the IRs are aligned *after* referencing, and with an acoustic timing reference they are: reconstructing the IRs from the exported 2026-08-10 sweeps puts every peak within **0.013 ms (4.5 mm)** of its own $t = 0$, and the L–R offset is **4.5 mm** — the two speakers are equidistant to the millimetre, exactly as the geometry says.

Timing matters in exactly one place: 3c forms a complex L/R average at each microphone position, before spatial phase is discarded. That is why the acoustic timing reference is required. Once the position sums exist, nothing downstream uses measured phase at all.

Why this is the only window you will ever set — three quotes from the Trace Arithmetic documentation:

- “*The currently applied impulse response window settings are used for each trace.*” → what you set on L and R is what every downstream operation consumes.
- “*The result uses the same window settings as trace A*” → derived trace arithmetic already consumes the windowed data. There is nothing to re-apply.
- “*Any frequency-dependent settings are excluded, applying an FDW to the result would amount to applying the window twice.*” → re-applying is not neutral, it is a second regularisation. REW deliberately clears the FDW flag on results. **Do not put it back.**

And one thing you cannot repair afterwards:

*“...or for division and inversion, **which use windows that span the entire resulting impulse response** as the result is typically not causal.”*

The divide’s output is deliberately un-windowed. **If the inputs were wrong, the output cannot be fixed.** This is why step 2 comes before everything.

[!] **The window is reversible. The traces built from it are not.**

Untick the FDW, press Apply Windows, and an original capture returns to exactly what it was. But trace arithmetic creates a **new measurement, frozen at the moment you press the button.** Change the window on L afterwards and LX, LX-MP, the division and the exported WAV all keep the old data — they do not recompute.

Set the window first, then do the arithmetic. If you change the window, redo the chain from step 3.

Step 3 — Reduce the captures to three divisors

You have ten windowed captures. The rest of the procedure needs exactly **three** traces out of them:

	what it is	used by
L-SP	the left channel, averaged over the five positions	the target, and the per-channel filter above 80 Hz
R-SP	the right channel, same	the target, and the per-channel filter above 80 Hz
LR-SP	the mono sum , averaged over the five positions	the common filter below 80 Hz (§11)

The word “family” Used throughout this step, and it means exactly one of three sets:

family	the five traces in it
L family	L C, L F20, L B20, L L20, L R20 — the left speaker measured at the five positions
R family	R C, R F20, R B20, R L20, R R20 — the right speaker at the same five
LR family	LR C, LR F20, LR B20, LR L20, LR R20 — built in 3c, one per position

A family is a set of five *positions*, always one channel or one sum — never a mixture of L and R. Each family gets averaged into exactly one spatial trace, and those three traces are the output of step 3.

The two rules Everything else in this step is bookkeeping.

Rule 1 — Phase is only meaningful *within* a position. Combine L with R while they still share a microphone point and a timing reference. Once you average across positions the phase is gone, and no later operation can recover it.

Rule 2 — Never change the L-to-R level ratio at a position. That ratio is not an error. It is the quantity that decides how deeply the two speakers cancel at that seat, which is the entire subject of §11. Any gain you apply at a position must be applied to **both** channels equally, or not at all.

Do the sub-steps in order. Trace arithmetic produces a **frozen** result: changing an input afterwards does not update anything already built from it.

3a — Collapse any repeat sweeps (optional) *Skip this if you took one sweep per channel per position.*

Repeat captures at one point measure repeatability. They are not extra spatial samples, and feeding all of them into the spatial average would silently weight that position more heavily than the others.

1. Select the repeats of **one channel at one position**.
2. Look at their impulse responses. Under an acoustic timing reference they overlay: **do not shift them merely because an alignment action exists**.
3. Choose **Vector average**. Name the result after the position — L C, overwriting the individual names.
4. Repeat for the other channel.

Vector average is right here because these are coherently timed measurements of the same source at the same point, so it also reduces uncorrelated noise.

[!] **Cross corr align has no role in this procedure**

It is the obvious button to reach for, and there are three places you might reach for it. All three are wrong.

Here, on the repeats. REW: *“time aligns the currently selected measurements by cross correlation of their windowed impulse responses, using the measurement which appears first in the list of those selected as the reference.”* Under an acoustic timing reference there is nothing to align. REW writes its own estimate into every measurement note — `grep -o 'Delay [-0-9.]* ms' *.txt` — and on the 2026-08-17 set the spread across repeats is **10 µs on L and 19 µs on R**, under 2° of phase at 225 Hz. If it is ever larger than a sample the timing reference did not hold; that is a measurement fault, and the fix is to sweep again.

In 3c, across the two channels. Never. This is aligning two different sources, and REW’s author is explicit that it does not work: > *“Cross correlation alignment is unlikely to work well if the sources are > different, e.g. if you were to try to align measurements of both left and > right speakers.”* — John Mulcahy

Worse, it would destroy the L-to-R arrival difference, which is the one phase quantity this procedure actually consumes. *(The same caution applies to the familiar sub-to-main alignment: that is also two different sources, and it is done by optimising the summed response, not by cross-correlating the two.)*

In 3d, across positions. Pointless. **RMS average** discards phase entirely, so time-aligning its inputs cannot change its output by so much as a hundredth of a decibel.

Cross correlation earns its place only ahead of a **Vector average** or an **RMS + phase avg.** of one source at several positions — a phase-correction branch this guide does not take. Mulcahy again: > *“Use cross correlation alignment and vector averaging if you need to work > with an impulse response, for EQ RMS or dB averages are often more > suitable.”*

3b — Level alignment: what it is for, and when to skip it REW’s only sentence on the subject:

*“If the measurements were made at different positions (spatial averaging) it is usually best to first use the **Align SPL...** feature to remove overall level differences due to different source distances.”*

and the action itself:

*“**Align SPL...**, which adjusts all the selected SPL traces so that they have the same average SPL over a selected span.”*

REW’s author, asked directly how much this matters:

“Aligning SPL is important if the measurements are from significantly different distances from the mic.” — John Mulcahy

“Significantly different” is the whole criterion, and it is a statement about your geometry, not about the software. The rest of this sub-step works out where the threshold falls for a cluster that maps one seat.

That is the whole of the official specification. It does not say what level the traces are aligned *to*, it does not say which traces you should put in the selection, and it provides no way to read back the offset it applied. So the rule below is derived here, from what the alignment is for, and then measured.

What it is for — weighting hygiene, not physics. RMS average is a power average, so a position that happens to be 3 dB louder carries twice the weight of the others. If that extra level is an artefact of where the microphone sat rather than something the correction should chase, it quietly biases the average toward that one seat’s shape. Align SPL exists to take that bias out.

It is **not** a correction of the response, and it is **not** applied to make channels match. The absolute level of the finished divisor is set later, by the target level in step 5.

At a 20 cm cluster, skip it. The geometry says why. Moving 20 cm along a 3.45 m path changes the direct level by $20 \cdot \log_{10}(3.65 / 3.45) = 0.49$ dB, and only for the forward and back points; the lateral points barely change their distance at all. Half a decibel of weighting error, spread across five positions in a power average, is not a thing worth correcting.

Measured, on the 2026-08-17 set — the finished common filter below 80 Hz, under four alignment policies, each compared against the exact one derived below:

policy	difference in the shipped filter
align nothing at all	0.01 dB rms, 0.04 dB max
Align SPL on each family separately	0.09 dB rms, 0.17 dB max
one offset per position, taken from L alone	0.06 dB rms, 0.21 dB max
L equalised against R at each position	0.39 dB rms, 0.68 dB max

Do nothing, and the filter is within 0.04 dB of perfectly aligned. For a cluster that maps one seat, this sub-step is optional. Skip to 3c.

It stops being optional when the cluster gets wide. Sample several seats a metre apart and the same arithmetic gives $20 \cdot \log_{10}(3.95 / 2.95) = 2.5$ dB of pure distance — five times the spread, and now worth removing.

If you do align: one offset per position, derived from both channels. Two constraints, and they fix the recipe completely.

It must not change the L-to-R ratio at any position (Rule 2). So the two channels at a point get the *same* number, and the sum built from them in 3c inherits it automatically.

It must be derived from both channels together — from the total energy arriving at that seat, $(|L|^2 + |R|^2) / 2$ averaged over the span. This is the part that is easy to get wrong. Move the microphone 20 cm to the left and the left speaker gets louder while the right gets quieter; the *seat* has not changed its overall distance from the pair at all. An offset computed from L alone would read that lateral asymmetry as a level error and take it out of both channels — pushing the already-quieter right channel further down. The power sum of the two is blind to the asymmetry and sees only what actually changed: how far the seat is from the pair as a whole.

The two disagree exactly where you would expect. On the 2026-08-17 set:

position	offset from L alone	from R alone	from both
C	+0.18 dB	−0.46 dB	−0.16 dB
8 cm left	+0.17	−0.37	−0.11
8 cm right	+0.12	−0.01	+0.07
10 cm forward	+0.12	−0.01	+0.07
33 cm back	−0.60	+0.85	+0.14

The single-channel columns swing by up to 1.45 dB against each other; the combined column stays inside a quarter of a decibel, which is what a cluster this size should produce.

In REW, since there is no way to read back what **Align SPL** decided: export the ten captures, take each one’s average level over 500 Hz–2 kHz, combine the pair at each position as a power average, and subtract the mean of the five. Then apply each result to both of that position’s captures with **SPL offset** → **Add to data**. If that is more bookkeeping than you want, use **Align SPL** on each family separately and accept the 0.17 dB in the table above — it is what REW documents, and it needs nothing read back.

Remember what this whole sub-step is worth here first. For a 20 cm cluster the table above already showed **doing nothing** (0.01 dB rms, 0.04 dB max) beats every alignment policy, including

Align SPL itself (0.09 dB rms, 0.17 dB max) — aligning is not fundamental to this procedure, it only earns its place once the cluster spans a metre or more. Everything below is for that wider-cluster case, or for whoever reaches for **Align** SPL anyway; if you already decided to skip this sub-step, skip this too and go to [3c](#).

If you do use it, the **Align** SPL... dialog’s two fields are worth setting rather than leaving alone, since it is already the less accurate of the two “do align” options in the table above and there is no reason to add band mismatch on top:

field	set to	why
Alignment center	1000 Hz	geometric centre of 500 Hz–2 kHz
Alignment Span	2 octaves	500 Hz → 2 kHz — the same band this section’s by-hand method already uses above

500 Hz–2 kHz is the same region the caution in step 1 already singles out as the one where microphone aiming matters and the capsule stops being effectively omnidirectional — comfortably above this room’s ≈ 166 Hz Schroeder frequency, so the level differences it measures reflect distance, not modal structure. (1000 Hz is $\sqrt{(500 \times 2000)}$; $500 \rightarrow 2000$ Hz is $\log_2(2000/500) = 2$ octaves.)

(This project’s own recollection of the factory-default centre/span — not independently confirmed, and REW’s published help documents the feature’s purpose but not its default field values — is 500 Hz, 2 octaves, i.e. 250 Hz–1 kHz: closer to the Schroeder region than the band above. Worth checking on a clean install if it matters to you; it does not change the conclusion that this whole sub-step is optional at this cluster size.)

[!] Never equalise L against R

The tempting move is to select **L L20** and **R L20** together and align them, so that at the 20 cm-left position the left trace comes down and the right comes up until they match. **Do not.**

At that seat the left speaker really is closer and really is louder. That difference is what the listener’s head experiences there, and it is the quantity that decides how completely the two speakers cancel at that point. Equalise the two and you compute a mono sum for a listener who is not in the room, and you fill in some of the cancellation that [§11](#) exists to protect.

On this set the two channels differ by only 0.8–1.3 dB over 500 Hz–2 kHz, so the damage is modest — the deepest null at each position moves by about 0.2 dB, and the shipped filter by 0.39 dB rms. It is still the worst of the four policies in the table, and it is the only one that is wrong in principle rather than merely imprecise. The error grows with any real imbalance: a channel gain mismatch, one speaker nearer a wall, a seat off the centre line.

Whatever you do, verify it: for each position, the L-minus-R level difference over 500 Hz–2 kHz must be exactly what it was before you touched anything. If it moved, an alignment reached across the two channels.

3c — Form the mono sum at each position This is the step Rule 1 exists for, and the only place inter-channel phase is ever used.

Do: for each of the five positions, select its **L** and **R** capture and choose **Vector average**. Name the results **LR C**, **LR F20**, **LR B20**, **LR L20**, **LR R20**.

Use Vector average, not Vector sum:

action	result	
Vector average	$(L + R) / 2$	✓ same SPL scale as one channel and as the target
Vector sum	$L + R$	× 6.0206 dB higher; would ask for 6 dB of spurious common cut

If you took the simultaneous **L+R** sweep at **C**, prefer it — it is the real acoustic sum, with no assumption that the room held still between two sequential sweeps.

A physical L+R sweep is $L + R$, so it sits **6.0206 dB above** REW's $(L + R) / 2$ and must be normalised before it can stand in for one:

1. make a response copy, so the raw capture stays untouched;
2. right-click the **SPL & Phase** graph → **SPL offset** → **-6.0206 dB** → **Add to data**;
3. name it **LR C** and use it in place of the calculated one. Keep the calculated version alongside as a check.

The 6.0206 dB is exact, and it has been verified twice

Second geometry. The 185 cm set contains a genuine measured L+R sweep beside its two channels. The complex sum of L0 and R0 reproduces the measured LR to **0.502 dB rms** over 20–250 Hz (median +0.010 dB), and the same measured sweep sits **+6.031 dB** above the vector average. Different room, different REW version, different hardware — same 6.02 dB.

This geometry. At the centre point, measured L+R sits a median **6.009 dB** above the calculated vector average over 20–225 Hz. After the exact **-6.0206 dB** normalisation the two agree to **0.276 dB rms** at 1/6 octave and 0.224 dB rms at native resolution above 80 Hz.

The one place they part company is a razor-thin cancellation near 49 Hz, where a fraction of a degree of sequential phase drift moves an almost perfect null between bins. That is an argument for preferring the measured sum, not against the substitution: swapping calculated for measured changes the finished five-position divisor by only **0.174 dB rms** below 80 Hz.

3d — Average across positions — three RMS averages Now, and only now, discard spatial phase.

Do: for each family in turn, select its five traces in the All SPL legend, right-click the graph and choose **RMS average**.

select these five	name the result
L C, L F20, L B20, L L20, L R20	L-SP
R C, R F20, R B20, R L20, R R20	R-SP
LR C, LR F20, LR B20, LR L20, LR R20	LR-SP

No further alignment here. If you applied position offsets in 3b the sums already carry them, and a family is never aligned a second time.

What you are giving up, deliberately. REW: *“Phase is not taken into account, measurements are treated as incoherent. . . the result has the magnitude data from the source measurement and no phase data.”* From here on the three divisors are magnitudes, and that is the intent — **R3** has the full ledger of which phase this procedure uses, which it discards, and which it deliberately refuses to correct.

Why RMS average and not the alternatives:

	why not
Vector average	position-dependent phase cancels and manufactures new nulls — the disease you are curing. REW: <i>“most appropriate for multiple measurements taken from the same position”</i>
RMS + phase avg.	identical magnitude, but it attaches a vector-averaged phase that means nothing across positions. REW: <i>“The resulting impulse response may have significant acausal content as the relationship between magnitude and phase that would normally hold is broken. As a result the average requires larger left windows than usual.”</i> Step 4 discards the phase anyway, so all you would buy is a longer left window
dB average	a level average, not a power average: <i>“gives equal weight to peaks and dips which masks the magnitude difference between them.”</i> REW allows it for deriving an EQ target from smoothed traces, but warns that <i>“with unsmoothed data the dips would have a disproportionate effect”</i> — and trace arithmetic here runs on unsmoothed data (§9)

All three families must contain the same five positions. Never build a five-position channel average against a four-position sum — the 80 Hz splice assumes both divisors describe the same set of listening points.

“Never average L with R” — and why LR-SP is not a violation of it

Standard REW practice is emphatic that the two channels stay in separate groups: average the five L into one trace, the five R into another, and never put a left and a right capture in the same selection. The reason is that **Align SPL** forces everything selected to a common level, so a mixed selection would erase the genuine left-to-right balance along with the positional scatter.

This procedure obeys that rule exactly. L-SP and R-SP are built from their own five captures and nothing else. Neither is ever contaminated by the other, and the target in step 5 is an average of the two *finished* channel traces, not of ten mixed captures.

LR-SP is a **third trace**, not a merged channel. It is never used as a channel response, never equalised against L-SP or R-SP, and never replaces either. It exists for one job: to be the divisor below 80 Hz.

And it exists because of the very risk the standard advice names. Correcting two channels independently is widely reported to produce a dip in the summed response wherever the channels’ phases diverge — which is §11’s failure, measured here at 2.5 dB of extra loss across 40–62.5 Hz. The usual escape is a single shared stereo filter, at the cost of never correcting either channel properly. The 80 Hz splice takes the shared filter where the problem is and the per-channel filters everywhere else. **The mono sum is the answer to that caveat, not an exception to the rule that produced it.**

3e — Bake the crossover correction into the channel averages

1. File → Import → Impulse Response → X801.wav. Name it X801 (revised). **Leave every window control on it alone** — see the warning below.
2. **Check its level, and correct it if needed.** REW’s impulse-response import has no calibration reference for a filter WAV, so it computes an SPL from the raw sample values rather than reporting the 0 dB an all-pass actually is — on this file, that reads as $\approx +117$ dB. Right-click the SPL & Phase graph → **SPL offset** → **−117 dB** → **Add to data** until the trace sits at 0 dB. Skip this and every trace built from X801 (revised) — LX, RX, and everything downstream through FLX/ FRX — inherits the same +117 dB, which the cut-only **Max gain** clamp in step 7 will then fight rather than the real correction curve.
3. Trace Arithmetic: $LX = A \times B$, $A = L\text{-SP}$, $B = X801$ (revised).
4. Likewise $RX = R\text{-SP} \times X801$ (revised).

Why: you are going to invert the system *as it will actually play*, and it will play through the crossover correction. Everything you then look at on screen — LX, RX and their average — is the real corrected system.

A useful thing to know: because X801 has magnitude 0.00000 dB everywhere, $|LX| = |L|$ exactly, and therefore LX-MP and L-MP are **identical**. This sub-step is magnitude-neutral, and the final filter would come out the same if you skipped it. Do it anyway, for two reasons: the traces you inspect are then honest about what you will hear, and if X801 is ever replaced by something that is *not* a pure all-pass (Xo801, which adds a bass-alignment term, is such a thing) the sub-step becomes load-bearing without warning.

[!] Multiply, *then* take minimum phase — never the other order

“Identical” above is not a coincidence of this particular file; it follows from what **Generate minimum phase** actually computes. Minimum phase is the Hilbert transform of **log-magnitude**, and the transform is linear:

$$\begin{aligned}\text{minphase}(L\text{-SP} \times X801) &= \text{Hilbert}(\log|L\text{-SP}| + \log|X801|) \\ &= \text{Hilbert}(\log|L\text{-SP}| + 0) && \text{since } |X801| = 1 \\ &= \text{Hilbert}(\log|L\text{-SP}|) = \text{minphase}(L\text{-SP})\end{aligned}$$

X801’s magnitude contributes exactly nothing to the Hilbert transform, so **as long as minimum-phase generation is the last operation**, it does not matter whether you multiply by X801 before or after — both routes land on $\text{minphase}(L\text{-SP})$, with none of X801’s actual phase θ_X in it. That is the mathematical content of “LX-MP and L-MP are identical.”

Reversing the order — minimum phase first, multiply last — is not the same operation, and it is wrong. L-MP already has its phase fixed at $\text{Hilbert}(\log|L\text{-SP}|)$. Multiplying it by X801 afterward *adds* X801's real, physical rotation θ_X on top:

L-MP	$\times$	X801:	magnitude	$ L-SP \times 1 = L-SP $	(same as above)
			phase	$Hilbert(\log L-SP) + \theta_X$	$\leftarrow$ X801's real phase, now present

That result is **not minimum phase** — θ_X is exactly the excess phase §2 defines a crossover all-pass to be, and a minimum-phase trace by definition carries none. Use it as the step 7 divisor and the division stops being the safe, magnitude-only inversion steps 4 and 8 exist to guarantee: R4’s own table is explicit that dividing by a non-**MP** trace “inverts magnitude **and** ... excess phase” — acausal, pre-ringing. The filter would then try to *invert* the crossover rotation it is about to receive cleanly and separately at step 9, defeating the orthogonality that section relies on (**LFilter** contributing “exactly nothing to crossover phase”).

The rule, stated once: minimum-phase generation must be the last thing that happens to a trace before it is used as a divisor. Multiplying by a flat-magnitude trace before that point is free — §above shows why — but multiplying afterward is not a stylistic variant, it is a different, and broken, result.

And why LR-SP is left alone. The same all-pass is applied to both channels, so it cancels out of their ratio: $(L \cdot X + R \cdot X) / 2 = X \cdot (L + R) / 2$, and the magnitude of the sum is unchanged. Since LR-SP is taken to minimum phase in step 4, which discards phase anyway, multiplying it by X801 would change nothing. Measured on the 2026-08-10 pair at 50 Hz: the vector average of LX and RX reads 67.603 dB, the vector average of the un-multiplied pair 67.602 dB. If X801 is ever replaced by something that is not magnitude-flat, build LR-SP from LX/RX-scale traces instead.

[!] Never window X801

It is an all-pass whose energy is spread symmetrically over ± 1365 ms by design. An FDW or a shortened right window would truncate the very phase rotation it exists to apply, silently degrading the crossover correction.

The FDW has already done its work on every original capture before L-SP and R-SP were formed. Do not reapply it to LX or RX. In step 9, FLX is built with **A = X801**, so it inherits X801's unwindowed state, which is right for a finished filter.

Step 4 — Minimum phase, first time

Do: on LX, use **Generate minimum phase** and name the new measurement LX-MP. Likewise make RX-MP from RX and LR-SP-MP from LR-SP.

(Naming note: *LR-SP* keeps its *-SP* tag through this step — *LR-SP-MP*, not *LR-MP* — because it goes straight from spatial average to minimum phase with nothing in between. *L-SP/R-SP* don't carry *-SP* into *LX/RX* because they pass through *3e's × X801* first, which is its own rename, not a dropped tag. Matches what both *../DRC-120.blue* and *../DRC-120.green* actually export.)

dialog option	set to	why
Cal file effects	included	you are modelling the acoustic response as measured, and the mic calibration is part of what the measurement means
LF tail	yes , at or just below the first measured bin: 16 Hz for the 2026-08-17 set , 15 Hz for the older set	required — without it the minimum-phase transform corrupts the magnitude it is supposed to preserve. See the callout below
Slope	24 dB/oct for a ported box (12 for sealed) — but 12 dB/oct regardless of enclosure if the LF-tail corner (the field just above, <i>not</i> the sweep-start frequency — REW floors it at sweep-start + 1 Hz) is $\geq 20/\sqrt{2} \approx 14.14$ Hz, i.e. within ½ octave of this project’s 20 Hz correction-band low edge	matches the speaker’s physical roll-off; halves the group delay near the corner when the two sit close. See “the corner’s distance from the band edge” below

dialog option	set to	why
HF tail	no	measured: the error above 1 kHz is 0.000–0.002 dB. The traces run to the top of the sweep — 22.05 kHz on the 2026-08-17 set, 24 kHz where the sweep went to Nyquist — which is six octaves above the correction band, far enough that the edge cannot reach it. Confirm with the 6a subtraction rather than assuming

[!] **The LF tail is not optional — this guide said “no” and was wrong**

A minimum-phase copy has exactly one obligation: **leave $|H|$ unchanged**, and supply the phase the magnitude implies. With **No LF tail** REW does not satisfy it. Measured on the 2026-08-12 build, max $|\text{error}|$ in dB, where every entry should be zero:

conversion	5–15 Hz	15–25 Hz	25–40 Hz	40–225 Hz
LX → LX-MP	13.41	0.58	0.03	0.33
RX → RX-MP	8.88	0.53	0.03	0.04
LR → LR-SP-MP	11.15	0.56	0.02	0.03
FL → LFilter	0.21	2.41	0.30	0.13

The mechanism. Minimum phase is obtained by a Hilbert transform of the log magnitude, which is a *global* operation — it needs the response defined across the whole spectrum. Truncate it and the transform rings against the edge, decaying away from it. That is why the error is largest at the bottom of the data and falls to ~0.03 dB by 25 Hz, and why the top end is spotless: the traces run to Nyquist, so there is no upper edge at all.

Why the corner goes at 15 Hz, not the default 20. Real sweep data starts at **15.01 Hz**; everything below is already synthetic, inherited from X801’s wider span through the trace multiply. A corner at 15 Hz replaces only synthetic data. A corner at 20 Hz would extrapolate over 5 Hz of genuine measurement — and that sits inside the band-limit blend (14.1–28.3 Hz), so it reaches the filter. If the field will not go below 20, take 20; the cost is small because the filter is fading toward unity through there anyway.

Put the corner at or below the first measured point, not just above it. The 2026-08-12 rebuild used 16 Hz against data starting at 15.0146 Hz, and the ~1 Hz of overlap left a **single-bin –140 dB null at exactly 15.01 Hz** in the exported WAV, where the tail splices onto the measurement. It is harmless — one bin, below the correction band, in a region already 29 dB down, and R8 does not see it — but it is free to avoid. Set the corner to the sweep’s start frequency and the splice has nothing to disagree with.

Slope: make the tail resemble what the trace physically is. For a measurement, the speaker’s own roll-off — **24 dB/oct** for a ported box, 12 for a sealed one. But see the next paragraph: if the corner sits near the band edge, the physical slope can cost more group delay than the gate allows.

Not 0 dB/oct, even though the dialog offers it and step 8 recommends it there. The two steps tail different objects. Step 8’s FL/FR are filters that are already, genuinely unity below the band — 0 dB/oct states a true fact about them. LX/RX/LR-SP here are *measurements* of a loudspeaker that really does roll off below its bass extension; 0 dB/oct would assert it keeps outputting at full level arbitrarily far below that, which is false. And because minimum phase is a *global* Hilbert transform of log-magnitude, an untrue magnitude story below the corner buys an untrue phase story inside the correction band, not just below it — the corner-distance discussion just below is that same mechanism at work for a slope that’s merely too steep, not wrong in kind.

Nor does the 24→12 dB/oct fallback license going further, down to 0 dB/oct, even to kill the group-delay bleed-through more completely. It would: a flat continuation induces no transition-related group delay at all, so in isolation it beats 12 dB/oct on that one number. But 24 and 12 dB/oct are both models of a real enclosure — a real box could plausibly be either alignment — so choosing between them trades one believable physical story for another, in a region cut-only clamps out of the final filter regardless of which you pick. 0 dB/oct isn’t a shallower version of that same trade: **no physical driver outputs flat to DC**, so it stops being a model of the loudspeaker

at all and becomes a known-false assumption chosen because it makes a test number smaller — the same mistake as §5’s cautionary tale, from the opposite direction: that one trusted an artifact because the numbers said so; this one would substitute a fabrication for missing data for the same reason. **Stop at 12 dB/oct.**

The corner’s distance from the band edge is a group-delay budget, not just a splice-quality one. A 24 dB/oct tail is a 4th-order high-pass; its group delay peaks just above the corner and is still large half an octave up. REW floors the corner field at **sweep-start + 1 Hz**, so a sweep that started at 16 Hz forces the corner to 17 Hz — only ½ octave below a 20 Hz band edge. **“Within ½ octave of 20 Hz” means the corner sits at or above $20 \div \sqrt{2} \approx 14.14$ Hz** — half an octave is a factor of $\sqrt{2}$, not 2, so this triggers more easily than it sounds: a 12–13 Hz sweep-derived corner clears it, a 15–16 Hz one (this project’s actual 120.blue/120.green captures, which start ≈ 15 Hz rather than step 1’s recommended 12–14) does not. The band-limited division then inherits that phase and the group-delay acceptance test fails at ~ 21 Hz: measured **+22 ms against the 10 ms gate**, with the filter’s *magnitude* already clamped flat there, so it is purely the tail. Moving the **target’s** LF cutoff (step 5) does not touch this — it is a different corner. Three ways out, cheapest first: - **Drop the tail to 12 dB/oct.** Halves the group delay near the corner. The gentler magnitude roll below the corner only asks for more sub-corner LF, which cut-only clamps — so it costs nothing real. Apply it in **both** minimum-phase steps (here and step 8). - **Raise the common-filter low edge to 25–28 Hz** (step 7). REW centres the one-octave blend on the limit, so a 25 Hz edge is still ~ 75 % faded in at 21 Hz; 28 Hz clears it. Costs ~ 1 dB of deep-bass taming. - **Re-sweep from ≤ 12 Hz** so the corner lands ≥ 1 octave below the band edge. Not 10 Hz on a ported woofer — see step 1.

Verify, do not assume. Export the copy and its source with **Smoothing: None** and subtract. $|LX-MP| - |LX|$ must sit at the ~ 0.03 dB level across the whole correction band. This is a two-minute check that catches a class of error nothing downstream can undo.

Why this step exists — it is the most important conceptual step in the chain. You are about to divide by this trace. What you divide by determines what gets inverted:

divisor	what the filter would try to undo	verdict
LX (the raw measurement)	magnitude and the room’s excess phase — every reflection, at one point	acausal, pre-rings, valid at one microphone position
LX-MP	magnitude, and exactly the phase that the magnitude implies	✓ causal, minimum delay, robust

Dividing by the raw measurement is the single most common way to produce a filter that measures beautifully at the microphone and sounds wrong. The minimum-phase copy throws away the part of the response that is not correctable, before it can reach the division.

Step 5 — Build the target

Do:

1. Select LX and RX in the All SPL legend, right-click the graph and choose **RMS average**. Name the result **L-R RMS average**.
2. Open it in the **EQ window** and save its shape as the target: **Target L-R RMS average**. Adjust the house curve here if you want one.

Why the L/R average and not each channel’s own response: it is already a partial spatial average — a null in the left path is filled by the right — and it gives both channels a **common** target, so correction does not pull the stereo image. Targeting each channel at its own smoothed response instead would let each chase its own private nulls, which is exactly the failure in §5.

Why go through the EQ window: what comes out is a **smooth curve**, not a copy of the measurement. It is also where the house curve goes. Note that a target shape is **magnitude only**, which is fine: only its magnitude survives step 8.

Why not RMS + phase avg.: its magnitude is identical, but its phase is a vector average and the resulting impulse may carry significant acausal content. A target is magnitude-only, so attaching a phase to it serves no purpose. Use the plain **RMS average**.

The target level — let REW calculate it, then check it this way The house curve is a **shape**, defined as 0 dB at 200 Hz. The **target level** is what anchors that shape in absolute SPL, and it is the single number that decides whether the corrected bass joins the rest of the spectrum or sits above or below it.

Press Calculate and let REW set it. REW anchors the target to the speaker’s **midrange reference level**, which is the right choice and is not obvious. Measured on the 2026-08-10 LR.rms+phavg:

region	level
correction band, 20–225 Hz	74.07 dB
300 Hz – 3 kHz (midrange)	68.80 dB
REW’s calculated target level	69.20 dB

REW’s number lands within **0.4 dB of the midrange**, and the correction band runs **5.3 dB hot** against it. That elevation is exactly what the filter is there to remove.

The check that matters: does the band edge join what is above it?

The filter is unity above 225 Hz, so whatever the correction leaves at the top of the band must match the untouched response just above it — otherwise you build a step at 225 Hz. With the calculated level:

180 Hz	200 Hz	225 Hz	225–315 Hz (untouched)
68.4	68.7	68.9 dB	68.5 dB

Continuous to a few tenths. **Do not** set the target level by matching the measurement’s own in-band level at 225 Hz — that keeps the room’s bass elevation instead of removing it, and leaves the whole band ~4 dB above the midrange. If REW’s calculated level looks “too low”, this is why: the number to compare it against is the **midrange**, not the neighbouring bass.

The target’s LF cutoff — move it to 5–10 Hz, below the correction band The target shape carries a **low-frequency cutoff**, and REW’s default puts it at **20 Hz, 24 dB/oct** — right at the bottom edge of the match range. Move the **frequency** to 5 or 10 Hz. **Leave the slope at REW’s default** — see below for why, unlike step 4, it isn’t a choice worth making here at all.

Frequency is the only knob that matters — not slope, not “match the speaker”

It is tempting to import step 4’s rule (match the physical roll-off) here. Don’t — it’s the wrong mechanism entirely, for two independent reasons:

1. **The target isn’t a measurement of anything.** It’s a *want*, built by hand in the EQ window, not the response of a real loudspeaker. There is no physical roll-off for it to be faithful to.
2. **The target is never minimum-phase-transformed.** Step 4’s slope choice was load-bearing because **Generate minimum phase** is a *global* Hilbert transform: an unphysical magnitude assumption below the corner distorts the phase it computes everywhere, including inside the passband. The target has no such step — “*a target shape is magnitude only... only its magnitude survives step 8*” — it is used directly as operand A in a division restricted to the 20–225 Hz match range. There is no mechanism by which its slope could propagate anywhere.

So the only thing that matters is **geometry**: get the corner far enough below 20 Hz that its bend doesn’t reach into the match range, which the table below shows the *frequency* alone accomplishes. Once the corner sits at 5–10 Hz, the transition is clear of 20 Hz regardless of how steep it is — there is nothing left for a slope choice to do.

Does that corrupt the calculated target level, given the speakers produce nothing at 10 Hz? No — and the reason is worth being clear about, because the worry is reasonable. REW fits the target to the measurement over the **match range**, 20–225 Hz. What the target does *below* 20 Hz is outside the fit entirely and never enters the calculation. What a 20 Hz cutoff does instead is bend the target **inside** the match range:

target attenuation from the LF cutoff	20 Hz	22 Hz	25 Hz	31.5 Hz	40 Hz
cutoff at 20 Hz	−3.01	−1.66	−0.67	−0.11	−0.02 dB
cutoff at 10 Hz	−0.02	−0.01	−0.00	−0.00	−0.00 dB
cutoff at 5 Hz	−0.00	−0.00	−0.00	−0.00	−0.00 dB

With the cutoff at 20 Hz, **10.9 % of the match range is pulled down by more than 0.5 dB**; at 10 Hz or 5 Hz, none of it is. Averaged across 20–225 Hz the depression is **−0.185 dB** at $f_c = 20$ and **−0.001 dB** at $f_c = 10$, so moving the cutoff shifts REW’s calculated target level by about **0.2 dB** — the target no longer has to sit slightly high to compensate for its own bend. That is the whole effect on the level, and it is negligible.

It does not make the filter reach below 20 Hz either. The division is band-limited to 20–225 Hz and reverts to unity below, blended over one octave; and max gain 0 dB clamps every boost regardless. Extending the cutoff changes **the target’s shape, not the filter’s reach** — there is no excursion risk from a target that is defined at 10 Hz.

Why this should clear two of the three acceptance-test failures

On the 2026-08-11 build, two of the three [R8](#) failures localise to **20.32 Hz** (narrowest feature) and **19.59 Hz** (group delay) — both sitting on the band edge, not on anything acoustic.

A 24 dB/oct corner *inside* the correction band is a steep feature, and the division reproduces it faithfully in the filter. Steep means high Q, and high Q means simultaneously a narrow feature and a long group delay — which is precisely the pair of tests that fail, at precisely the corner frequency. It is the same trade as §4: you cannot have the sharp edge and the short filter.

Putting the corner at 5–10 Hz leaves the target smooth across the whole 20–225 Hz band, so the filter has no steep band-edge feature to build and both tests should clear. The third failure is a separate matter and this will not touch it. **Verify on the rebuild rather than assuming** — this is a prediction from the mechanism, not a measured result.

Choosing a target the filter can actually reach This is the decision that determines whether the correction does anything at all, so it is worth more than a moment. With max gain at 0 dB the filter only cuts, therefore:

$$\text{achieved} = \min(\text{target}, \text{measurement})$$

A target that sits **above** the measurement produces no correction whatever — not a partial one, none. So the target is not a wish. It is a **ceiling you slide down under the measured curve**, and the measurement is its upper bound.

Measured here after step 2 (FDW 12 cycles), in dB relative to each channel’s own level at 200 Hz — i.e. exactly that upper bound:

	20	25	31.5	40	50	63	80	100	125	160	200 Hz
L	+3.7	+6.1	+1.9	−5.2	+6.9	+0.1	+6.0	+2.7	+7.2	+1.0	−0.1
R	+4.6	+7.8	+2.8	−0.8	+2.0	+5.0	+11.2	+0.3	+9.4	+6.5	0.0

This room is not bass-shy at 120 cm. Below 160 Hz it runs *above* its own 200 Hz level almost everywhere. What it actually has is **lumps at 80 and 125 Hz** and a **hole near 40 Hz** — and the hole is 4.4 dB deeper on L than on R, so it is one channel’s path, not a room property, and must not be filled (§6 rule 2). The job here is mostly **cutting**, not lifting.

[!] Ask for more lift than the room offers and you get no correction at all

The deployed chain used a **+6.5 dB** shelf. Against the table above that target sits over the measurement across most of the low band, so the clamp engaged and the filter did nothing:

	25 Hz	35 Hz	45 Hz	63 Hz
correction produced	0.00 dB	+0.01	+0.01	+0.01

Roughly half the correction band received no correction. A house curve you cannot reach is not a gentle preference — it silently switches the filter off.

Four candidate shapes, and what each actually delivers Computed as $\min(\text{target}, \text{measurement})$ on the real FDW-12 data, both channels, everything in dB relative to 200 Hz:

target	deep 20–45	upper 63–160	balance	ripple	band cut	deepest cut
<i>no correction (the room)</i>	+2.8	+6.1	−3.3	4.21	—	—
A as deployed, +6.5 shelf	+2.6	+1.6	+1.0	2.59	57 %	8.6 dB
B reachable, +3.0 shelf	+1.1	+0.3	+0.8	1.84	68 %	10.0 dB
C +4 low, −3 upper-bass	+1.5	−2.2	+3.7	2.07	78 %	12.9 dB
D flat	−0.8	−0.2	−0.6	1.41	74 %	10.6 dB

“**Balance**” is the number that decides how much bass you hear: deep minus upper, i.e. how far the bottom two octaves stand clear of the upper bass. Read the first row and the problem is obvious — **this room is upper-bass dominant by 3.3 dB**. That is the 80 and 125 Hz lumps, and they mask everything below them.

- **D (flat) is the wrong answer.** It cuts the deep bass along with everything else and ends up with *less* low-end weight than doing nothing.
- **A (as deployed) is doing least of all** — 57 % of the band, because most of the shelf is unreachable and clamps.
- **C is the one to use if you want low end.** It gives **+3.7 dB of balance against A’s +1.0** — nearly 3 dB more apparent weight — with *less* ripple than A, and it costs only 1.1 dB of absolute deep bass relative to A.

More bass comes from a deeper cut, not a taller shelf

With a cut-only filter a taller shelf means *less* cutting, so it buys loudness and lumpiness together. The lever that actually delivers low end is the **upper-bass scoop**: pulling 80–160 Hz down un-masks the 20–45 Hz region that is already there. Cutting is always reachable, so this direction is free in a way that lifting never is.

Recommended: shape C — about **+4 dB below ~45 Hz**, through 0 dB near 70 Hz, to a broad **−2 to −3 dB centred around 110–120 Hz**, back to 0 dB by 200 Hz. Start the scoop at −2 dB and deepen it by ear; too much hollows out male voice and cello. The extra cut costs headroom (12.9 dB at the worst peak, against 8.6 dB now), which is amplifier gain, not excursion.

The delivered scoop runs deeper than the drawn one — draw shallower than you want to hear

$\text{achieved} = \min(\text{target}, \text{measurement})$ is only exact against the **minimum-phase copy** the division actually uses (step 4). The real, non-minimum-phase measurement can and does sit *below* that copy, and that residual is not corrected — it is exactly what step 4 discards before the division so the filter doesn’t chase it.

Measured on the 2026-08-25 **Rscreen** build (target drawn at **−1.74 dB**, 90–140 Hz):

	target	delivered L+R	gap
90–140 Hz	−1.55 dB	−2.88 dB	−1.33 dB

Pointwise on the R channel it is worse — −8.0 dB at 90 Hz and −9.0 dB at 140 Hz against a −1 to −2 dB target, because R’s non-minimum-phase content is larger there than L’s. A scoop drawn at −1.74 dB, aimed at “start at −2 dB,” was delivered at −2.9 dB — already past the “too much” the guidance above warns about, and audible as thinness in the presence region rather than as “more balance.”

Draw the scoop shallower than the target you actually want, by roughly the size of this gap for your own measurement — check it the same way, by comparing `LX8-MP + Fcommon x Fdiff` (what the filter believes it's hitting) against the real corrected prediction (`LR.filtered` or equivalent). Do not assume the gap is 1.3 dB; that number came from one FDW-8 build and one room. Verify it before trusting it, and re-check after any change to the FDW window or the min-phase reconstruction, since both change how much non-minimum-phase content step 4 throws away.

If the scoop keeps landing thin: don't use one

The scoop is the aggressive option and easy to over-shoot — the delivered gap above, plus a room that only needs a decibel or two of upper-bass trim, can leave 90–160 Hz gouged and the presence region audibly thin. The conservative alternative is a **no-scoop Harman / Olive shape**: a smooth LF shelf ($\sim +4$ to $+5$ dB), flat mids, no dip. It gives up some of C's un-masking “balance” for a fuller, more forgiving low mid, and it is the right choice when the spatially-averaged room is already close to neutral (deep-vs-upper within ~ 2 dB). `tools/housecurve/housecurve.py` writes it as `tools/housecurve/house-curve-harman.txt` (+4) and `tools/housecurve/house-curve-harman-fuller.txt` (+5); the 120.blue multi-point build uses the +5. The file carries a -1 dB/oct treble tilt for REW's target-level calc only — the filter is unity above 225 Hz, so the tilt is not delivered.

Two things the target cannot do:

- **Anything above 225 Hz.** The filter is unity there (step 7), so the target's shape above the band is decoration. If you want a treble tilt it has to be a separate broad shelving filter, not this target.
- **Fill the 40 Hz hole.** It is destructive interference on one channel. Boosting it costs excursion and does not fill it.

(If you want more low-end than the reachable ceiling allows, the honest fixes are placement and treatment, not the target — the same conclusion as §5's.)

Step 6 — Stop and verify

Do: export LX-MP with **Smoothing: None** and run **two** checks.

6a. The minimum-phase copy must have preserved the magnitude. Export LX too, subtract, and require $|LX-MP| - |LX|$ to sit at the **~ 0.03 dB** level across 20–225 Hz. Repeat for RX-MP and LR-SP-MP.

This is a property, not a tolerance: a minimum-phase copy changes phase and nothing else, so any visible deviation is an artifact. If it fails, the LF tail in step 4 is off — see the callout there. Nothing downstream can undo an error in the divisor, because everything downstream divides by it.

6b. No peak below 200 Hz may be both narrower than ~ 30 FFT bins (≈ 11 Hz) and more than ~ 6 dB above its surroundings.

Why peaks and not dips. Because max gain 0 dB (step 7) already disarms the dips, and only the peaks reach the WAV:

in the divisor	what the division makes of it	survives?
narrow dip	filter wants a narrow <i>boost</i> $\rightarrow$ clamped to unity	no
narrow peak	filter builds a narrow <i>deep cut</i>	yes

This matters, because an honest divisor routinely contains narrow dips and they are not a defect. On the 2026-08-11 build:

trace, un-smoothed, 20–200 Hz	narrowest peak	narrowest dip
LX-MP (the divisor)	14 bins, 3.9 dB @ 35.5 Hz	6 bins, 40.1 dB @ 188 Hz
RX-MP (the divisor)	23 bins, 4.1 dB @ 50.5 Hz	13 bins, 4.1 dB @ 57.9 Hz
FL (what ships)	38 bins, 5.8 dB @ 98.9 Hz	12 bins, 3.0 dB @ 35.5 Hz

A rule applied to *any* feature would have rejected LX-MP over a 6-bin dip and thrown away a build whose filter is comfortably inside the threshold. The same applies to the 74 Hz front-wall null that a 12-cycle FDW exposes on the left channel (§8): across 60–90 Hz LX-MP swings **35.7 dB**, while FL swings **11.7 dB** — at the null itself the filter sits at exactly 0.00 dB, unity, with the cut confined to the shoulders. **Windowing revealing a null is working as intended; it is not a Step 6 failure.**

Why the FDW still matters even so. The clamp defuses dips, not peaks, and un-windowed modal peaks are exaggerated by reverberant energy that belongs to the room and not to the direct sound — invert those and you cut deep narrow notches fitted to one microphone point. It also keeps the shoulders around a null from being over-cut right next to a unity notch. §5’s disaster was a dip, and cut-only alone would have blunted it; the FDW is what keeps the rest of the filter honest.

If a narrow tall peak does survive, the FDW in step 2 did not take. Common causes: it was applied to the wrong measurement; Apply Windows was never pressed; or the arithmetic in step 3 was done *before* step 2 and is holding stale data. Thirty bins is the acceptance threshold from R8, applied one step early — and the surest version of this test is to run it on FL itself once step 7 is done.

*(Optional belt-and-braces, step 4a: bake the smoothing in by round trip — apply 1/6 octave to LX/RX, export as text **with that smoothing selected in the export dialog**, re-import, and take the minimum phase of the re-imported traces. Redundant if step 2 was set properly. Doing neither is what produced §5.)*

Step 7 — The division

Two divisions and a multiply, not one division. One shared setup (7a, 7b) and then two operations per channel (7c, 7d). The reason is §11: below 80 Hz the two speakers cancel each other at the seat, and dividing by each channel separately deepens that cancellation.

7a — take the spatial sum built in step 3. Use LR-SP-MP.

Do not vector-average LX and RX here. L-SP and R-SP are spatial RMS averages: their position phase is already gone, so a vector average of them is not a mono sum of anything. The sum must be formed L+R at each position (3c) and RMS-averaged across positions (3d), in that order.

7b — the common filter, below 80 Hz. Trace Arithmetic, **A over B**:

field	value
A	Target L-R RMS average
B	LR-SP-MP
Lower / upper frequency limit	20 Hz / 80 Hz
Max gain	selected, value 0.0 dB

Name it **Fcommon**. There is only one, shared by both channels.

If the 20–200 Hz group-delay gate fails near the low edge

The **20 Hz** lower limit above is the default and correct starting point. But it sits close to [step 4](#)’s minimum-phase LF-tail corner — see that step’s “corner’s distance from the band edge” callout for the mechanism. If `tools/drc_acceptance.py` fails the group-delay excursion test just above 20 Hz, the fix that touches only this step is: **raise this lower limit to 25 Hz** (band-limit blend is centred on the limit, so 21 Hz only fades in ~75 %, but it trims the excursion). Measured on the 2026-08-31 multi-point rebuild: 27 → 22 ms at 21 Hz — an improvement, not on its own a pass. Combine with the other two fixes in step 4’s callout if 25 Hz alone doesn’t clear the gate; the trade is ~1 dB less taming right at 20 Hz.

7c — the per-channel filters, above 80 Hz. Trace Arithmetic, **A over B**:

field	value
A	Target L-R RMS average — <i>the same A as 7b</i>
B	LX-MP (the minimum-phase copy, not LX)
Lower / upper frequency limit	80 Hz / 225 Hz
Max gain	selected, value 0.0 dB

Name it `Fper_L`. Repeat with `B = RX-MP` for `Fper_R`.

7d — combine. Trace Arithmetic, **A times B**: `Fcommon × Fper_L → FL`. Same with `Fper_R → FR`.

Both factors are cut-only, so the product is cut-only: **Max gain remains selected and set to 0.0 dB for both divisions.** The 80 Hz splice is seamless because REW's band-limit ramps are complementary — §11 has the algebra and the measured agreement.

(To reproduce the older single-division behaviour — per-channel everywhere — skip 7a and 7c, and in 7b use `B = LX-MP` with limits 20 Hz / 225 Hz. The tables below were measured on that version, and the boost/cut asymmetry they show is unchanged.)

Why the frequency limits are what makes this a filter, not a curve. Outside the limits the result reverts to **unity gain**, blended over one octave centred on the limit. Measured on the deployed file:

10 Hz	20 Hz	225 Hz	250 Hz	300 Hz	1 kHz	15 kHz
+0.04 dB	+0.02	−3.26	−0.25	+0.01	+0.01	+0.01

Unity to within 0.04 dB outside the band, and the one-octave blend is visible in the 225 → 250 → 300 Hz run. Without the limits you would be exporting the whole target curve as a filter.

What max gain 0 dB actually does: the filter only ever cuts

Read off the deployed chain — the correction the target asked for, against what the division produced:

	25 Hz	35 Hz	45 Hz	63 Hz	80 Hz	125 Hz
Target —	+2.2	+4.3	+2.6	+2.5	−2.0	−4.4 dB
LX-MP (wanted)						
what came out	0.00	+0.01	+0.01	+0.01	−3.7	−4.8 dB

Every boost is clamped flat; every cut passes. That is the setting doing exactly what it says, and it is the right conservative choice — a deep null at one microphone point is destructive interference, not a deficiency of output, and filling it with gain wastes excursion and does not fill it.

Know what it implies: **the correction is cut-only.** It will lower the peaks toward the target and leave the dips where they are. If your target looks nothing like the result, this is why.

(Figures read off the 96 ppo exports, which carry display smoothing; the divide itself ran on unsmoothed data, so treat the decibels as indicative and the boost/cut asymmetry as exact.)

Why the beta's Max gain control is selected. V5.40 beta no longer has the old division Regularisation percentage. It has a **Max gain** check box and value. Selecting it does two jobs: it clamps boost, and with frequency limits it makes the result revert to unity outside the selected band. With the value at 0.0 dB the correction is cut-only. This is not merely an argument from principle; it is what the deployed filter measures:

FLX-trimmed-48k.wav, 20 – 225 Hz	
maximum	+1.19 dB
minimum	−7.24 dB

The filter never boosts. If broad boost is ever deliberately allowed, raise the **Max gain** value by the exact amount intended. Do not deselect the check box: in current beta that also changes the outside-band behaviour of a limited division from unity to the offset numerator. For orientation, the depths that exist in this data, measured relative to the divisor's own average level over 20–225 Hz:

divisor	deepest dip below its average
L, raw (no FDW)	−30.4 dB
L, FDW 20 cycles	−24.1 dB
L, FDW 12 cycles	−13.7 dB
R, FDW 12 cycles	−8.5 dB

After step 2 there is no raw canyon left to feed to the division — **the FDW has already done that job**. Max gain = 0.0 dB then refuses all remaining requests to boost a null.

[!] **The feature that caused 1.3 seconds of ringing is 3.5 dB deep**

FLX-trimmed-48k.wav at native 0.3662 Hz resolution, around the null. These are absolute filter gains, not relative to anything:

28.198 Hz	28.564	28.931	29.297
+0.87 dB	−1.64	−2.35 dB	+1.19
GD +71 ms	−101	−105 ms	+81

The bottom of the dip sits at **−2.35 dB**. What matters is the **excursion across it**: from +1.19 dB in the bin next door down to −2.35 dB is **3.54 dB**, and it happens over **two bins** — **0.73 Hz**. That is the whole feature. It carries **105 ms of group delay** and rings for **1348 ms**.

Depth is not what hurts. Width is. A 3.5 dB wiggle you would never notice on an amplitude plot is a Q-40 resonator because $Q = 28.93 \div 0.73 \approx 40$, and §4’s law converts 0.73 Hz directly into 1.4 seconds. This is the whole reason step 2 exists, and the reason no guard in this step can substitute for it.

*(How that dip got there, since the clamp had pinned the magnitude flat: LX-MP-INV+0dB is magnitude-flat to −0.03 dB at 28.93 Hz but its **phase** jumps from +32° to −116° in one step. The clamp constrains magnitude and leaves the phase discontinuity behind; step 8 then re-derives phase from a full-resolution magnitude that was never really flat. A hard clamp constrains amplitude, not bandwidth* — that part of the old diagnosis stands.)**

Step 8 — Minimum phase, second time

Do: Generate minimum phase on FL → LFilter. Likewise FR → RFilter.

dialog option	set to	why
Cal file effects	not included	a filter has no microphone. Including the mic calibration would bake the microphone’s response into what you play as step 4: without it the transform corrupts the magnitude. This is the copy where it does the most damage
LF tail	yes — see the note below	
Slope	shallowest offered — the beta dialog goes down to 0 dB/oct — or match step 4’s choice if step 4 dropped to 12 dB/oct for the group-delay budget	FL is a filter, not a loudspeaker — it has no physical roll-off to match. See “Slope here is a real choice” below
HF tail	no	as step 4

This is where the LF-tail error hurts most

Step 4’s copies are damaged at 5–15 Hz, mostly *below* the correction band. Here the damage lands **inside** it — FL is unity below ~14 Hz, so the transform’s sharpest edge is the band-limit transition itself, and that is where it rings:

$\backslash \backslash \text{LFilter} \backslash \backslash$ − $\backslash \backslash \text{FL} \backslash \backslash$	5–15 Hz	15–20 Hz	20–25 Hz	25–40 Hz	40–225 Hz
max error	0.21	2.41	2.39	0.30	0.13 dB

In the exported WAV this resolves into a **−11.4 dB notch at 19.96 Hz**, identical on both channels to 0.01 dB — deterministic processing, not acoustics. It is 0.23 Hz wide at −6 dB, narrower than REW’s 0.366 Hz display grid, so **the plot shows only a −4 dB dip and cannot resolve the core**. It is what fails the [R8](#) narrowest-feature and group-delay tests, both of which land at ~20 Hz.

Slope here is a real choice, not just conditioning. FL is a filter, not a loudspeaker — it has no physical roll-off, and it is already unity below the band limit. A steep tail imposes a subsonic high-pass on the deliverable. That may be welcome, but decide it rather than inherit it: pick the **shallowest slope offered** if you want the filter left as designed.

0 dB/oct is on the table, and it is the correct choice for this exact goal. The beta dialog’s slope field goes all the way down to 0, and 0 dB/oct means the minimum-phase transform extrapolates the magnitude below the corner as **flat, at unity**, rather than tilting it into any roll-off at all — matching what FL already *is* there (§ above: unity below the band limit) instead of imposing a shape on top of it. Where 24 or 12 dB/oct asks “how steep should the imposed high-pass be,” 0 dB/oct is the answer “there should be no imposed high-pass” — the softest filter this field can produce, and the one that changes FL least. Used on the current `../DRC-120.green` build (`L.Filter/R.Filter`, per its audit), with no ill effect traceable to it: the build’s remaining acceptance failures sit at 78–81 Hz, an octave-plus above where this slope acts.

Do not confuse 0 dB/oct with No LF tail. They are opposites, not degrees of the same thing. No LF tail leaves the magnitude undefined below the corner and is the documented failure two paragraphs up — 2–13 dB of corrupted magnitude, the mistake this whole callout exists to retract. 0 dB/oct is the LF tail **switched on**, with the mildest slope it can extrapolate with. Turning the tail on is not optional; choosing 0 dB/oct for its slope is.

The target is unaffected, because Step 5 builds it from LX/RX rather than the minimum-phase copies (that is the reason for the instruction). If you are rebuilding after fixing this, `LRrms+phavg`, the target shape and the calculated target level all stand — redo from Step 4 onward only.

Why a second time, when the divisor was already minimum phase. In exact arithmetic the quotient of a magnitude-only numerator and a minimum-phase divisor is minimum phase, so this ought to be a no-op. In practice it is not, for three reasons: the max-gain clamp and the one-octave band blend both perturb the phase; the division “uses windows that span the entire resulting impulse response as the result is typically not causal”; and the target has zero phase, which is not the same as having none. Measured on the deployed chain the step moved the 20 Hz phase from -41.6° to -2.3° .

What it guarantees: a causal, minimum-delay filter whose phase is exactly the Hilbert transform of its magnitude. That is the definition of the thing you meant to build.

Step 9 — Bake the crossover correction in, last

Do: Trace Arithmetic, **A times B:**

- `FLX = A = X801 (revised)`, `B = LFilter`
- `FRX = A = X801 (revised)`, `B = RFilter`

Why last, and why `A = X801`: keeping the room correction (`LFilter`) and the crossover correction (`X801`) as separate traces until the end means you can re-derive one without re-baking the other, and you can A/B the crossover correction by exporting with and without this multiply. And with `X801` as trace A, the result inherits `X801`’s window settings — i.e. none — which is correct for a finished filter.

Why the cascade is safe: `LFilter` is magnitude-derived minimum phase and contributes exactly nothing to crossover phase (§3); `X801` has exactly flat magnitude and contributes exactly nothing to the magnitude correction. They are orthogonal. There is no double-correction to worry about.

Step 10 — Export

Do:

1. Trim IR to windows on `FLX / FRX` if you want to set the latency.
2. Export impulse response as WAV: **48 kHz, 32-bit float**. REW writes **131072 samples** — there is no tap count to set, the export length is fixed.
3. Name them `FLX-trimmed-48k.wav / FRX-trimmed-48k.wav`.
4. Alongside the WAVs, export the same two traces as **text**, **Smothing: None** — `FLX-trimmed.txt / FRX-trimmed.txt`. Deployment doesn’t convolve with these; it uses them as an independent statement of what the filter is supposed to do, and checks the WAVs against it. Two of the ten files [step 12](#) needs — the rest of that bundle isn’t produced yet at this point.

Trimming never changes the filter’s response. FLX and FLX-trimmed were compared bin by bin: identical to **0.0000 dB rms at every frequency**. All it does is move the impulse within the file — it sets latency, nothing else.

So **trim every filter you build**, including every rebuild: each pass through the chain produces a new FLX/FRX and each one gets trimmed and exported as normal. What trimming is not, is a **remedy**. If step 11 fails, re-trimming the same filter cannot change any of the three test results, because none of them depends on where the impulse sits in the file. The defect is upstream and only step 2 can reach it.

There is no “number of taps” in REW. rePhase has one; REW does not. The equivalent lever is the IR window of step 2, which is why that step matters so much. If you need a genuinely shorter filter, window the exported WAV outside REW with a **tapered** (Hann/Tukey) window — never a hard truncation, which produces ripple.

Step 11 — Accept or reject, before deploying

[!] Do not verify a five-position filter at one position

Remeasuring with the filter in place is the right final check, but measure at **three or more of the five points**, not just the centre. A multi-position filter matches the target at the centre *worse* than a single-point filter would — that is the trade you bought, not a defect (§6).

On the phase trace expect: the crossover rotation gone (X801, step 9), a minimum-phase rotation consistent with the magnitude changes, and **the room’s reflections untouched**. If reflections look corrected, something divided by a raw measurement instead of a -MP trace.

Do:

```
./tools/drc_acceptance.py ../DRC-120.blue/FLX-trimmed-48k.wav \
    ../DRC-120.blue/FRX-trimmed-48k.wav --plot check.png
```

Exit status 0 = pass. Full description in [R8](#).

If it fails, do not deploy and do not post-process. Go back to step 2, lower the FDW cycles, and re-run the chain from step 3. A filter that fails these tests fails audibly.

Step 12 — Prepare the export bundle for open-media-drc

This is the condensed version. `../open-media-drc/doc/FILTER_PROVENANCE_AND_RESPONSE.md` and the open-media-drc provenance documentation is the reference — read it for the declare/tag/build/publish chain that comes after this step. This step is only about getting the **export directory** into the shape that chain, or the web UI’s live installer, requires. Nothing below opens BruteFIR or touches a deployed room; it is still REW work.

The ten files, and what each one actually is (per open-media-drc/scripts/README.md; case-insensitive, .txt optional):

file	what it is	source
L.txt, R.txt	the room before correction, left/right	a sweep — see the warning below
LR.txt	the room before correction, the pair — REW’s vector average	a sweep, Vector average of the two above
FLX-trimmed.txt, FRX-trimmed.txt	the finished filter’s own response	step 10.4
FLX-trimmed-48k.wav, FRX-trimmed-48k.wav	the deployable impulses	step 10
L.filtered.txt, R.filtered.txt	the room, with the filter applied — the prediction	built here
LR.filtered.txt	the filtered pair	built here, Vector average of the two above

Pick **one** aggregate convention and hold it: LR throughout, never mixed with L+R.txt/L+R.filtered.txt. the open-media-drc provenance documentation confirms this explicitly — “for a filter built by this guide’s procedure

it is *vector_average*: step 3c forms the vector $L + R$ and then subtracts 6.0206 dB, which is $(L + R) / 2$.” A design that mixes LR with L+R anything is refused outright, because the two curves would then be measuring different things.

[!] `L.txt/R.txt/LR.txt` must be a sweep — not the traces this

procedure has been building all along

This is the mistake the checklist below exists to prevent. L/R as they exist by this point in the session carry **step 2’s FDW** — that is correct and required for building the filter, and exactly wrong for these three files. open-media-drc’s own classification is explicit: `L.txt/R.txt/LR.txt` are “*the room before correction... a sweep*” — a direct, unregularised capture, not a divisor.

The open-media-drc install will not catch this for you. Its unsmoothed check looks for one specific thing — REW’s **Smoothing** header, the 1/3-octave-etc. display-smoothing setting — and an FDW-windowed export states **Smoothing: None** just as honestly as a raw one does. FDW is a different REW feature entirely, and nothing in the validator inspects it. An FDW-windowed `L.txt` would be **accepted, published, and shown on the page as the measured room** — quietly wrong, exactly the failure mode step 2 itself warns about (“the window sits in the wrong place... quietly, with no error message”), one stage further downstream than that warning covers.

Do:

1. **Get unwindowed copies of the Main Listening Position’s L/R. L C/R C** (or whichever capture is the primary seat) already carry the FDW applied in step 2 — do not touch those; the filter still needs them exactly as they are. Instead, **duplicate** each one in REW’s measurement list, and on the duplicates only: open **IR Windows**, untick **Add FDW**, press **Apply Windows**. Per step 2’s own note, this is fully reversible on an *original* capture — the duplicate returns to exactly what the sweep measured. Name the duplicates so the provenance is obvious in the session — this project’s own declared designs use an **.orig** suffix in the REW trace name (e.g. `L 120.Rscreen.orig`) for exactly this reason; the trace name and the exported file name don’t have to match.
2. **Set REW’s Smoothing menu to None** before exporting anything in this step. This is the separate control §9 already established doesn’t reach trace arithmetic — it *does* reach a plain text export, which is exactly the surface this step uses.
3. **Form LR.** Select the two unwindowed duplicates, **Vector average** — same operation as 3c, on different inputs. Export all three: `L.txt`, `R.txt`, `LR.txt`.
4. **Build the prediction.** Import `FLX-trimmed-48k.wav` and `FRX-trimmed-48k.wav` as measurements (File → Import → Impulse Response), the same mechanism step 3e used for `X801`. **Check the import level before trusting it** — REW’s impulse import has no calibration reference for a filter WAV in general, not only for an all-pass one; step 3e measured `X801` importing at +117 dB for exactly this reason. Verify against a point you know (e.g. the filter’s own passband gain from `FLX-trimmed.txt`) and correct with **SPL offset** → **Add to data** if it’s off, rather than assuming it imported at 0 dB.
5. **Multiply.** Trace Arithmetic, **A times B**: the *unwindowed* `L` duplicate × the imported `FLX-trimmed` → `L.Filtered`. Same for `R × FRX-trimmed` → `R.Filtered`. **Vector average** the two → `LR.Filtered`. Export all three, **Smoothing: None** again.
6. **Verify the set** before handing it to `new_filter_design.py` or the web installer: exactly ten files, one aggregate convention throughout, every text export’s header literally reads *** Smoothing: None** (not merely silent on the subject — an export that states no smoothing at all is refused too, because it can’t be *shown* to be unsmoothed), and none reaches past 24 kHz. These are the same checks the tool runs; this is catching them yourself first, with the file names in front of you rather than in a refusal message.

```
./tools/drc_export_preflight.py ../DRC-120.green/120.green.multipoint.txts
```

Run it on the export directory. It mirrors `new_filter_design.py`’s own file-role resolver — including the LR/L+R conflict check above — its unsmoothed/24 kHz checks, and the listening-position geometry comments the web UI’s room diagram reads out of L/R/the aggregate’s Notes field (front wall to MLP, MLP to speakers, speakers to front wall, floor-marker colour — case-insensitive, decimal commas normalised; REW writes the leading *** Note:** itself on export, so type only the sentence, e.g. **4.18m from front wall**).

Missing comments only warn, matching the real tool; conflicting values between files fail, also matching it. None of this needs SoX or a Git checkout of `open-media-drc` to run. It does **not** replace the TXT↔WAV residual check (that needs SoX + NumPy); a clean pass here still means running `new_filter_design.py --dry-run` before deploying, not skipping it.

7. **Commit.** The export directory must be a Git work tree, and all ten files plus the `.mdat` must be committed — not merely staged — before deployment can read them back later. See the current `open-media-drc` provenance documentation for what comes next: declare the roles, tag, build, publish. comes next: declare the roles, tag, build, publish.

Part IV — Reference

R1. The IR window dialog, field by field

Defaults as REW sets them on a full-range sweep:

field	default	what it does	change it?
Left window shape	Rectangular	taper applied before the reference time	no
Left width	500 ms	how much pre-arrival is kept	no — it does not affect resolution
Ref Time	~0	position of the window reference relative to $t = 0$	no — it is the IR peak, which is what the acoustic timing reference bought you
Right window shape	Rectangular	taper applied after the reference time	only on route B below
Right width	1000 ms	how much decay is kept	only on route B below
Add FDW	off, 15 cycles	the frequency-dependent Gaussian window	✓ ticked, 12 cycles

Frequency resolution = 1 / (right width). Settled by experiment.

REW’s help says the printed resolution corresponds to “*the current total window duration (left and right combined)*”. The dialog does not behave that way. Measured:

left	right	printed resolution	1/(L+R)	1/R
500 ms	1000 ms	1.00 Hz	0.67 Hz ×	1.00 Hz ✓
500 ms	500 ms	2.00 Hz	1.00 Hz ×	2.00 Hz ✓

The left width does not affect frequency resolution. It is not a resolution control; set the right width and leave the left alone.

Ref Time is not something you compute. REW puts $t = 0$ at the impulse peak, so Ref Time comes out at ~0 (0.042 ms on `L.120.Blue` — the residual offset of the peak from the sample grid). Move it only with the Impulse graph’s *Set $t = 0$ at cursor* / *Offset $t = 0$* actions.

Route B — a fixed window instead of the FDW

Only if you want a flat Δf across the band. This gives the **wrong shape** — a constant Δf is far heavier regularisation at 4 kHz than at 50 Hz — so prefer the FDW unless you have a specific reason.

control	set to
Right shape	Tukey 0.5 — change this <i>before</i> the width
Right width	173 ms → prints 5.8 Hz $\approx 1/6$ octave at 50 Hz (300 ms → 3.3 Hz is a gentler first try)
Left shape / width	leave
Add FDW	off

total window	Δf shown	$\approx$ octave fraction at 50 Hz
1.5 s (default)	0.67 Hz	1/52 — effectively none
500 ms	2.00 Hz	1/17
300 ms	3.33 Hz	1/10
173 ms	5.78 Hz	1/6
86 ms	11.63 Hz	1/3

Do not do both. You would be regularising twice, and the printed resolution would no longer describe the data.

R2. Choosing a window shape

Only relevant on route B. Read it anyway — it explains why route A avoids the question.

(a) six shapes as REW applies them. (b) the same six in frequency — **this is the smoothing kernel** the response gets convolved with. (c) each applied to the real L.120.Blue impulse response at a 173 ms window.

[!] Textbook window figures do not apply to a half window

The familiar numbers (Hann -31 dB sidelobes, Blackman-Harris -92 dB) are for **symmetric** windows. REW's right window is truncated at its own peak, so the signal meets a step at the reference time however smooth the taper is, and that step dominates the kernel. Measured on the actual half-windows at 173 ms:

shape	-3 dB width	leakage outside main lobe	kernel
Rectangular	5.13 Hz	9.23 %	10 nulls — oscillates
Tukey 0.25	5.86 Hz	6.90 %	7 nulls — oscillates
Tukey 0.5	6.59 Hz	5.66 %	3 nulls — smooth
Hann	8.06 Hz	9.67 %	0 nulls — smooth
Hamming	6.59 Hz	11.72 %	8 nulls — oscillates
Blackman-Harris	9.52 Hz	12.75 %	0 nulls — smooth

Blackman-Harris is the worst of both worlds here — widest main lobe *and* highest leakage, the exact opposite of its symmetric-window reputation.

Use **Tukey 0.5**, chosen on measurement rather than reputation: lowest leakage of the six and a smooth non-oscillating kernel, for only $1.28\times$ the resolution of rectangular. The distinction that matters is not sidelobe **level** — all six sit at -13 to -16 dB — but sidelobe **structure**. A rectangular kernel *rings*; a tapered kernel decays smoothly. Ringing in the kernel is ringing convolved into every feature of the response.

Why a time window is a smoothing kernel

Multiplication in time is convolution in frequency. Windowing the impulse response by $w(t)$ replaces $H(f)$ with $H(f) * W(f)$:

$$Y(f) = \int H(\nu) \cdot W(f - \nu) d\nu$$

Every output frequency becomes a **weighted average** of the input, with W supplying the weights — which is the definition of smoothing. A weighting function inside an integral operator is called its **kernel** (German *Kern*, “core”), and here **the kernel is literally $W(f) = \text{FFT}\{w(t)\}$** , which is what panel (b) plots. A window's shape in time and its kernel in frequency are one object seen from two sides, which is why “choose a window length” and “choose a smoothing bandwidth” are one decision.

window in time	kernel in frequency	behaves like
rectangular, length T	sinc, main lobe $\sim 1/T$, -13 dB sidelobes	crude smoothing at $\Delta f = 1/T$, plus ringing
Tukey / Hann, length T	wider main lobe, sidelobes $\ll$	clean smoothing at $\Delta f \approx 1/T$
Gaussian FDW, N cycles	Gaussian of width f/N	fractional-octave smoothing, $1/N$ per octave

One difference that matters. The smoothing menu convolves the **magnitude** — it averages $|H|$, so it can only make a null *shallower*. Windowing convolves the **complex** response, so contributions can cancel and a




Figure 4: Window shapes in time, their frequency kernels, and the effect on the real measurement

window can make a null *deeper*. REW’s help says the FDW has “an effect similar to applying a smoothing of the same octave fraction” — *similar*, not identical, and this is where it differs.

R3. Minimum phase — the two places you take it, and the one place you must not

The phase ledger

Four kinds of phase pass through this procedure and each is treated differently. Confusing them is the most common way to misread the chain.

phase	fate	why
across positions	discarded by the RMS average in 3d	it describes one seat’s interference pattern; correcting it optimises a point nobody’s head stays in
L against R at one position	used once, in 3c, then gone	it is the mono cancellation that §11 exists to correct. This is the only phase the procedure consumes, and why the acoustic timing reference is mandatory
the room’s excess phase	never corrected	inverting it needs an acausal filter, pre-rings, and is valid at one point (§2)
the crossover all-pass	corrected — by X801, in step 9	position-independent, deterministic, known in closed form; and invisible to a magnitude-derived inverse

So discarding measured phase and regenerating minimum phase is not a loss you absorb. **It is the mechanism that stops the filter attempting the second and third rows.**

One consequence worth knowing: because L-SP is magnitude-only and X801 is magnitude-flat, LX-MP is identical to the minimum-phase copy of L-SP itself. X801 does no work at 3e in a multi-position build; it earns its keep at step 9, where it is baked into the shipped filter.

#	applied to	producing	why
1	LX, RX, LR-SP	LX-MP, RX-MP, LR-SP-MP	every divisor must be minimum phase. Divide by a raw measurement and the filter tries to invert the room’s excess phase: acausal, pre-ringing, valid at one microphone point
2	FL, FR	LFilter, RFilter	the filter must be causal. The clamp, the band blend and REW’s un-windowed division output all leave residual non-minimum-phase content
×	X801	—	never. Its magnitude is 0.00000 dB, so its minimum-phase copy is a unit impulse — you would delete the filter entirely

That last row is measurable, not rhetorical:

minimum-phase copy of X801.wav:

```

peak                1.000000  at sample 0
energy anywhere but sample 0  2.96e-14 %
```

Cal file effects: *included* at step 4 (you are modelling the acoustic response as measured, and the mic calibration is part of what that means), *not included* at step 8 (a filter has no microphone). Below 200 Hz a decent measurement mic is nearly flat, so this is a small effect — but the reasoning is asymmetric and worth following.

R4. How to invert: $\div$, $1/A$ and $1/|A|$

REW offers several operations that all look like “invert”. They are not interchangeable.

operation	result	outside its frequency limits	guards it offers	use it?
$A \div B$, $A = \text{Target}$, $B = \text{LR-SP-MP}$	Target $\div$ the spatial mono sum	unity , blended over one octave when Max gain is selected	Max gain	✓ the common filter , 20–80 Hz (§11)
$A \div B$, $A = \text{Target}$, $B = \text{LX-MP}$	Target $\div$ one channel	unity , blended over one octave when Max gain is selected	Max gain	✓ the per-channel filter , 80–225 Hz
$A \div B$, $A = \text{Target}$, $B = \text{LX-MP}$	Target $\div$ measurement	unity , blended over one octave when Max gain is selected	Max gain	the older single-division form; deepens the 45–56 Hz mono cancellation
1/A on LX-MP	flat at a chosen level	unity , blended over one octave	Max gain , target level, exclude notches	only if you want a flat target and no house curve
1/ A on LX-MP	magnitude inverse	—	—	× produces a linear-phase result — symmetric pre-ring, and no phase correction at all
$A \div B$ with $B = \text{LX}$	inverts magnitude and room excess phase			× acausal, pre-rings, one-point-only

Why $A \div B$ and not 1/A: the division takes a **target curve**, so the house curve of step 5 comes along for free. 1/A takes only a scalar *target level*, so it can only aim at flat. If you are happy with flat-at-the-average across 20–225 Hz, 1/LX-MP in one step is perfectly respectable and gives you the **exclude notches** option, which $A \div B$ does not have.

The one thing 1/A has that division does not is that notch-exclusion checkbox, which drops notch-like features out of the inversion entirely rather than merely limiting their boost. If a null survives step 2, that is the one remaining lever — at the cost of the target curve.

Never 1/|A|. REW’s help: “*The 1/|A| and 1/|B| operations produce a linear phase result.*” A linear-phase filter is symmetric: half its energy arrives **before** the impulse. That is the opposite of what step 8 is for.

R5. The two boost guards in V5.40 beta

guard	belongs to	mechanism	effect on bandwidth	set it to
Max gain	division, inversion	optional hard clamp at a level; for limited division it also selects unity outside the limits	none — a clamp is a discontinuity, and a discontinuity is a narrow feature	selected, 0.0 dB — gives a cut-only, band-limited filter
Exclude notches	inversion only	drops notch-like features from the inverse	removes them entirely	off

Both limit boost. Neither limits bandwidth. So neither can prevent the failure in §5, which was a *dip* — 0.73 Hz wide, 3.5 dB from rim to bottom — in a filter that never boosted at all. **The FDW is not a backstop to these controls; these controls are a footnote to the FDW.**

With 12 cycles applied at step 2 the 38.3 dB null is already an 11.7 dB dip and the deepest point in the whole 20–225 Hz divisor is 13.7 dB below its average. The 0.0 dB clamp then refuses to fill it.

The obsolete **Regularisation** percentage belonged to older REW versions. In V5.40 beta 25 and later division uses a **Max gain** figure, and gain limiting can be disabled by deselecting its check box. Do not translate old percentage recipes into the current UI; set the desired maximum gain directly.

R6. Getting the room’s decay figure

You need one number to choose the FDW cycles: how long this room actually decays in the 40–100 Hz region. **No single method gives it below ~150 Hz.**

tool	gives	trustworthy where	use it for
RT60 (T20 / T30 / EDT / Topt)	numbers per band	above Schroeder only (~166 Hz)	mid/high-band RT, treatment work
Waterfall	3D ridges	qualitative	<i>which</i> modes ring; before/after
Spectrogram	2D map, colour = level	qualitative, $\pm 2\times$	<i>how long</i> — the referee

tool	gives	trustworthy where	use it for
Peak bandwidth	numbers per mode	below Schroeder, if cross-checked	the figure you actually use

Route 1 — from peak bandwidth. Since $T60 = 2.2 \cdot Q / f_0$ and $Q = f_0 / \Delta f$, the frequency cancels:

$$T60 = 2.2 / \Delta f, \text{ where } \Delta f \text{ is the } -3 \text{ dB bandwidth of the peak, in Hz.}$$

Decay depends *only* on bandwidth. A 5 Hz-wide peak decays in 440 ms whatever its frequency. Measure Δf off the **unsmoothed** trace with the cursor.

Its blind spot: a single-point magnitude measurement **cannot distinguish a genuine high-Q mode from a narrow interference peak**. Both are narrow, and the formula converts either into a decay time. On this dataset it returns 2215 ms at 53.8 Hz — physically impossible in 58 m³ — which is the §5 trap seen from the frequency side.

Route 2 — RT60 graph. Standard, but meaningful only **above** Schroeder. Below it the field is modal rather than diffuse, “reverberation time” is not well defined, and the analysis band filter has its own floor: a 1/3-octave filter at 25 Hz rings 205 ms unaided. Numbers from this route below ~150 Hz are not trustworthy.

Route 3 — spectrogram, as the referee. It cannot give a number, but it answers the one question the others cannot: *how long does energy at this frequency actually persist?* Here, the 40–100 Hz band fades in roughly **300–400 ms**.

Why no view can give you a number — the same law, again. Both waterfall and spectrogram are built on an analysis window of length W , and that window sets *both* axes: frequency resolution $\approx 1/W$, time smearing $\approx W$. To resolve a 5 Hz-wide mode at 50 Hz you need $W \geq 200$ ms; but a 200 ms window cannot resolve a decay much shorter than 200 ms, and the decay you are measuring is ~400 ms. There is no setting that gives both. $\Delta f \cdot \Delta t \approx 1$ (§4), now limiting the instrument you would use to check the instrument.

These views are reliable to a **factor of about two** — enough to reject 2200 ms, nowhere near enough to separate 300 from 400 ms. Referee, not instrument.

Reading a spectrogram without fooling yourself:

- Analysis window $\approx$ **300 ms** for bass work: long enough to separate modes ~3 Hz apart, short enough not to invent decay.
- Level range so the floor sits **35–40 dB** below the peak. Too wide and everything looks like it rings forever; too narrow and real decay vanishes.
- Read decay as *where a streak fades relative to where it started*, with the cursor — not by eye off a screenshot.
- **Ignore everything below ~20–25 Hz.** The sweep has little energy there and the window is many cycles long; the smear is the analysis, not the room.

So: route 1 for the numbers, route 3 to reject the impossible ones. Any peak whose bandwidth implies a decay far longer than the spectrogram shows is interference, not a mode, and **must not be corrected**. Cross-checked that way, the credible decays here are **300–400 ms** in the 40–100 Hz band, which is the table in §8.

Applying the same arithmetic to the features the September filter built:

feature in FLX	Q	source resonance it implies	verdict
28.93 Hz	40.4	T60 = 3.07 s	impossible
51.27 Hz	62.7	T60 = 2.69 s	impossible
81.67 Hz	11.9	321 ms	plausible
116.46 Hz	15.7	297 ms	plausible

Independent proof that the two worst features were never room modes — and the answer to “does the window throw away something real”: for exactly those features, no.

R7. X801 — what it is, and do not redesign it

X801.wav linearises the two Nautilus 801 crossovers and nothing else. Measured directly from the file:

	value
magnitude, 5 Hz – 23 kHz	–0.00000 to +0.00000 dB — a perfect all-pass
group delay fit, low crossover	375.4 Hz, Q 0.706
group delay fit, high crossover	3945.5 Hz, Q 0.686
residual of the two-all-pass fit, 25 Hz – 15 kHz	0.0076 ms rms
fitting the high crossover alone	26× worse
relative group delay, 20 Hz vs treble	–1.30 ms
length	131072 samples, 32-bit float, 48 kHz, centred

It contributes **nothing** to magnitude and therefore nothing to low-frequency ringing. It is ready to use as-is.

[!] X801.rephase on disk does not match X801.wav

The saved rePhase project has an **empty filter type** on the 355 Hz row, but the WAV measurably corrects both crossovers. **Open that project and re-export and you will silently lose the 350 Hz correction — which carries ~95 % of the group delay.**

The WAV is the artifact of record. If it must ever be rebuilt: low row at ~355–375 Hz, LR 24 dB/oct ($Q \approx 0.707$), which is what the fit says is actually in there.

How it is designed in rePhase: the *filter linearization* tab, entering the two crossover frequencies and their topology. That is the whole design, and it is the correct minimal one.

Not to be confused with Xo801.wav

DRC-185/Xo801.wav is the *same two crossovers* — they agree above 350 Hz to **0.0051 ms rms** — **plus** an extra low-frequency term: a 4th-order high-pass at **20.66 Hz**, i.e. a bass-reflex box alignment. It is not wrong, it is a different **scope**.

Use the crossover-only X801. Xo801's cost is genuinely small (1.6 ms of pre-ring at –40 dB), but its 20.7 Hz figure is a model fitted to the filter itself and never checked against a nearfield measurement of a real 801 — and below 100 Hz the room contributes several times more phase than the ~22 ms it corrects. Low cost, unverified benefit. Restore it only after measuring the speaker's rolloff nearfield.

Minor open point: the N801 crossover is nominally third-order, whose sum is an all-pass of $Q = 1.0$, while the fit says $Q \approx 0.706$ (LR4-like). If the real acoustic sum is $Q = 1.0$ at 350 Hz, about 0.6 ms of rotation remains uncorrected — at or under the Blauert & Laws audibility threshold. A refinement, not a defect.

R8. Acceptance tests

`tools/drc_acceptance.py`, in this directory. Run it on the WAV BruteFIR actually loads. Exit status 0 = pass.

```
./tools/drc_acceptance.py ../DRC-120.blue/FLX-trimmed-48k.wav \
    ../DRC-120.blue/FRX-trimmed-48k.wav --plot check.png
```

test	pass	September FLX
sharpest feature below 200 Hz	Q of 12 or less	Q 73 ×
group delay excursion, 20–200 Hz	< 10 ms	80 ms ×
gated-tone tail, any tone below 200 Hz	< 100 ms	1348 ms ×
magnitude, DC ... 20 Hz	flat, no correction	ok

[!] The first test used to be in FFT bins. That was wrong

Bin spacing is f_s / n , so a bin-based threshold **passes or fails the same filter depending on how long the exported file is**. Trimming one of these filters from 262144 to 131072 taps changed its verdict without changing the filter: the band-edge feature measures **9.07 Hz** wide in one and

9.37 Hz in the other — the same feature — but 49.6 bins against 25.6, because each bin doubled in width.

The criterion is now **Q = centre frequency ÷ bandwidth**, which is dimensionless and length-independent. The threshold follows from §8: an N-cycle FDW caps fractional bandwidth at 1/N, so a divisor windowed at 8–12 cycles cannot legitimately yield a feature sharper than about 12.

Calibrated against this project’s history:

	width	at	Q	
Sept 2025 deployed filter	0.30 Hz	28.9 Hz	73	× rang 1348 ms
no-LF-tail notch	0.27 Hz	20.0 Hz	73	×
74 Hz SBIR shoulder, 12-cycle build	5.78 Hz	73.2 Hz	12.7	× marginal
8-cycle build	13.1 Hz	99.6 Hz	7.6	✓
20 Hz band-edge transition	9.37 Hz	25.6 Hz	2.7	✓

Note this also changes *which* feature is reported: “sharpest” now means largest $f_0/\text{bandwidth}$, not smallest bandwidth. A 13 Hz feature at 100 Hz is sharper in the sense that matters than a 9 Hz one at 25 Hz.

Bins are still printed, for reference only.

The script uses a **matched-latency control** for the gated-tone test and applies a **4-period floor** — a 40 dB decay of a 28.7 Hz tone cannot resolve faster than ~139 ms, so demanding less would measure the envelope estimator rather than the filter. Verified behaviour: **X801.wav passes** all three; **FLX-trimmed-48k.wav fails** all three.

What to run it on: **FLX-trimmed-48k.wav** / **FRX-trimmed-48k.wav**. Also run it on **LFilter/RFilter** if you want to separate room-correction ringing from anything the crossover filter contributes — **X801.wav** passes on its own, so any failure is the room part.

By hand, if you prefer, import the filter as a measurement (**File** → **Import** → **Impulse Response**) and look at **Group Delay** first. It is one graph and it caught every failure in this dataset:

	28.93 Hz	51.27 Hz	79 Hz	116.5 Hz	max 20–200 Hz
FLX (L)	– 105 ms	– 97 ms	–0.9	–35	+80 ms @ 29.3 Hz
FRX (R)	+0.7	+0.8	– 65 ms	–32	+52 ms @ 75.4 Hz

A correction filter should show group delay of a **few milliseconds** below 200 Hz. Tens of milliseconds means you have built a resonator.

Seeing all three tests in REW, and comparing two filters

Every number the script produces has a REW view behind it. Import the WAV with **File** → **Import** → **Impulse Response** and it becomes a measurement like any other.

REW view	what it shows	script test
Waterfall (CSD) or Spectrogram	how long each frequency keeps ringing	gated-tone tails
Group Delay	resonators, as GD peaks	group delay excursion
SPL, Smoothing: None	narrow features	narrowest feature
Step Response	pre-ringing and overshoot	pre-peak energy

Waterfall settings for a filter (not for a room — the defaults assume a room): time range ~**400 ms**, rise time **100 ms**, **40 dB** range. Anything still standing at the back of the plot is what the script calls a tail.

To compare two filters, import both and use the **Overlays** window. All **SPL** and **Group Delay** overlay directly, which is the fastest way to see what a rebuild changed. Waterfalls do not overlay — flip between them with the measurement selector.

Thresholds, so you can judge audibility rather than pass/fail

The script’s thresholds are deliberately conservative — they are there to catch construction faults, not to predict audibility. These are the numbers for the second question. Both are approximate and the literature disagrees by roughly a factor of two.

Group delay (Blauert & Laws; Møller):

20 Hz	50	100	200	500	1 kHz
~45 ms	~32	~24	~18	~12	~6 ms

Modal ringing: audibility for bass decay is usually placed around **200–300 ms** for normal program material.

Applied to this project’s filters:

	worst group delay	verdict	worst tail	verdict
April build	–104.6 ms @ 28.9 Hz	~2.5× over threshold	>2000 ms @ 28.7 Hz	far over
8-cycle build	+17.2 ms @ 31.9 Hz	under (~38 ms there)	114 ms @ 79 Hz	under

The 8-cycle build still **fails** the script on group delay — 17.2 ms against a 10 ms criterion — while sitting **below the audibility threshold** at the frequency where it occurs. That is the criterion doing its job: it is a build-quality gate, not a verdict on the sound. Read a failure as “look at this”, not as “this is audible”.

What the guards cost

Rebuilding the left filter from its own magnitude with LF regularisation:

gated note	as deployed	1/6 oct < 200 Hz	1/3 oct < 200 Hz
28.7 Hz	1348 ms	1	1
51.2 Hz	787	72	16
79 Hz	1209	72	34
81.6 Hz	511	46	35
116.5 Hz	264	45	30
145.5 Hz	465	30	9
narrowest feature < 200 Hz	0.8 bins	31.5 bins	126.7 bins
change to the correction curve	—	1.29 dB rms	1.52 dB rms

1.29 dB rms of correction accuracy buys the removal of 1.3 seconds of ringing. That is the trade this whole procedure makes, quantified.

How the gated-tone test works

Worth understanding, because it is the one test that maps directly onto something you can hear. For each of ten tones — 28.7, 40, 51.2, 63, 79, 100, 116.5, 128.2, 145.5, 180 Hz, the room’s modal frequencies — the script:

1. builds a **1-second sine**, with 5 ms raised-cosine ramps at both ends so that switching it off does not itself click;
2. appends **2 seconds of silence** — the note stops dead;
3. **convolves it with the filter**;
4. takes the **Hilbert envelope**, reads the steady-state level from the 300 ms before note-off, and times how long the envelope takes to fall **40 dB** below it.

That time is the **tail**. It tests **the filter alone**, not the filter in the room: whatever it rings is *added* to the room’s own decay.

The **control** column is a pure delay pushed through the identical chain. It should ring for zero time, so whatever it reports is the measurement’s own floor — the envelope estimator needs several cycles to resolve a 40 dB drop, which takes longer as frequency falls. Hence controls of ~124 ms at 28.7 Hz falling to 3 ms at 180 Hz. The limit is $\max(100\text{ ms}, 3 \times \text{control}, 4000/f_0\text{ ms})$, the last term being four periods.

[!] **A single tone length is not a measurement**

The gate closes at whatever phase the tone happens to be in, and the envelope of a low tone ripples at $2f_0$, so the 40 dB crossing lands on either side of a ripple lobe. On an **unchanged** filter, 28.7 Hz returned 80, 82, 149, 217, 259 or 85 ms depending on a 2 % change in tone length.

The script therefore reports the **median over nine lengths** and an **IQR spread column**, and flags any entry whose spread exceeds half its median as **noisy**. **Entries flagged noisy are not evidence in either direction** — read the spread before believing any tail.

This is not academic. The single-length version made the 8-cycle rebuild above look *worse* than the 12-cycle one, on two filters that are magnitude-identical at 28.7 Hz to 0.04 dB and therefore cannot differ there. The **control** was hit harder than the filter: it drew 0 ms at 28.7 Hz while reading 94 ms at 40 Hz. A pure delay’s floor must grow monotonically as frequency falls, so that zig-zag was the tell — and because the limit is $3 \times \text{control}$, a control that collapses to zero drops the 28.7 Hz limit from 373 ms to 139 and fails filters that pass.

If the controls in your output are not monotonic in frequency, do not trust the tails.

Worked example: R8 blinded by processing, then not

The 2026-08-12 rebuild is the clearest demonstration of what these tests are for. The same filter, built twice from identical measurements, differing only in the **LF tail** of step 4 and step 8:

test	No LF tail	LF tail at 16 Hz
narrowest feature < 200 Hz	$\times$ 1.49 bins @ 19.96 Hz	$\checkmark$ 31.56 bins @ 73.24 Hz
group delay, 20–200 Hz	$\times$ +257.4 ms @ 20.32 Hz	$\times$ +24.1 ms @ 74.52 Hz
gated tones failing	5 , worst >2000 ms @ 28.7 Hz	2 , worst 228 ms @ 63 Hz

Two things to take from it.

First, the frequency response barely moved — the delivered mono sum changed by +0.01 to +0.12 dB in every band. All three tests improved by an order of magnitude while the thing you would look at in REW stayed put. That is the entire argument for running time-domain acceptance tests at all: **the defect was invisible on the magnitude plot**, and REW’s own 0.366 Hz grid could not resolve it even in principle (§ step 8).

Second, look at where the failures moved to. Every remaining one sits at **63–92 Hz**, and the filter there reads:

68 Hz	71	73	74	76	79	85	92
–2.68	–0.88	–0.04	–0.12	–1.51	–6.22	–10.37	–8.09 dB

That unity plateau at 73–74 Hz between two cuts is the front-wall SBIR null being clamped to unity by max gain 0 — the mechanism [step 6](#) describes, with geometry predicting the null at 75.3 Hz. Ten decibels of swing in half an octave is a high-Q feature, and a high-Q feature is a long tail: 63 and 79 Hz are the **shoulders of that null**.

So R8 stopped reporting an artifact and started reporting the one real acoustic problem in the room. **A failing test that has moved onto a feature you can point at in the geometry is a different kind of failure from one sitting on a band edge.** The first is a decision about how hard to correct; the second is a bug.

R9. Does the source correlation invalidate §11?

§11 divides by the **vector sum**, which is the response you get when both speakers reproduce *the same waveform*. Real music below 80 Hz is not necessarily the same in both channels, so the question is fair: is the divisor describing a situation your recordings are not in?

Correlation, in one paragraph

Two speakers can only cancel each other if they are **playing the same thing**. The 45–56 Hz cancellation is destructive interference between the left and right arrivals — which requires the two arrivals to carry the same waveform. Feed the speakers unrelated signals and there is nothing to cancel; the two just add in energy, like two people talking at once. The measure of “same waveform” is the correlation **r** between the channels, band-limited to the region of interest:

r	the two channels contain	consequence at the seat
+1	the identical waveform	the speakers cancel fully — the measured null is real
0	unrelated waveforms	no cancellation at all — powers simply add
−1	the same waveform, inverted	cancellation on the mono sum is at its worst

Note r runs from **−1 to +1**, and negative values do occur in real recordings. Note also that correlation is **not** the same as balance: equal energy in both channels with $r \approx 1$ gives a solid centre image, while equal energy with $r \approx 0$ gives diffuse bass with no location at all.

What the room measurement assumes

A sine sweep is perfectly correlated by construction. So the vector sum shows the $r = 1$ case — **the deepest the null can ever be**. The physically correct divisor for source correlation r is

$$|H_L|^2 + |H_R|^2 + 2r \cdot \text{Re}(H_L \cdot H_R^*)$$

which is the vector sum at $r = 1$ and the power sum — essentially the RMS average — at $r = 0$.

What is actually on the discs

Measured over 71 albums of this collection, one track each, band-limited to 20–80 Hz:

	median r	energy-weighted mean	range
all tracks	+0.65	+0.46	−0.57 ... +0.99
tracks with >3 % of energy below 80 Hz	+0.50	+0.44	−0.57 ... +0.99

Fully spread — Mendelssohn organ discs at 0.93–0.99, Shostakovich’s 1st and 15th at **−0.57**. At $r \approx 0.45$ the correct divisor sits between the two extremes, and it differs from the vector sum by **2.9 dB rms across 45–56 Hz, worst 4.5 dB**, while agreeing to within 0.25 dB everywhere else.

And yet it changes nothing — because of the clamp

That 2.9 dB is an error in the **divisor**, not in the **filter**. Carried through to what the filter does:

Hz	target	divisor r=1	divisor r=0.45	cut applied, r=1	cut applied, r=0.45
45	71.5	67.2	67.9	0.00	0.00
50	71.1	67.7	70.4	0.00	0.00
56	70.6	63.9	68.6	0.00	0.00
71	69.3	71.3	71.7	−2.06	−2.45
80	68.6	74.9	75.3	−6.33	−6.75

Across 45–56 Hz the divisor lies **below** the target on either assumption — by 3.5 dB at 50 Hz with $r = 1$, still 0.7 dB with $r = 0.45$. Both ask for a **boost**, and **max gain 0 dB clamps every boost**. The clamp is engaged on 100 % of bins in that band with both divisors, so the filter is identical.

The only real difference is **≈0.4 dB more cut at 71–80 Hz** — below session repeatability, and not worth a rebuild.

The general lesson, which is the reason this appendix exists

An error in the divisor is only an error in the filter **where the filter is acting**. Below 80 Hz this design spends most of its time clamped, so discrepancies there are absorbed rather than propagated. The same reasoning retires the narrow dips in [step 6](#) and the deep null at 74 Hz. **Before acting on a divisor discrepancy, carry it through `min(Target − divisor, 0)` and see whether it survives**. It usually does not.

(Caveat worth keeping: one 90-second excerpt per album is decent breadth and thin depth, and r varies within a work as much as between works. The conclusion above is robust because the clamp absorbs the whole range of r , not because 0.45 is a precise figure.)

R10. Troubleshooting

symptom	cause	fix
Woofers keep moving after the music stops	narrow features in the correction curve → high-Q resonators	FDW at step 2. Verify at step 6
Group delay spike of tens of ms below 100 Hz	same	same
Step 6 fails — features still 1–2 bins wide	the FDW never reached the data	check it was applied to every original capture , that Apply Windows was pressed, and that step 3 was done after step 2
Filter is not flat above 300 Hz	the division’s frequency limits were not set	step 7 — set 20 Hz and 225 Hz
Filter has energy <i>before</i> its peak	$1/ A $ was used, which is linear phase; or the divisor was LX not LX-MP	steps 4 and 7
Crossover phase unchanged after correction	expected — a minimum-phase inversion cannot see an all-pass	step 9; see §3
Applying 1/6 octave smoothing changed nothing	trace arithmetic ignores the smoothing menu	§9 — use the FDW
Correction sounds right at the mic, wrong one seat over	single-point correction of position-specific features	more positions (§6), or fewer FDW cycles
Changed the window, the filter did not change	derived traces are snapshots, not live views	redo the chain from step 3
Deep null got <i>deeper</i> after the FDW	the FDW convolves the complex response; that null is early-arrival, not late	it is real — do not try to fill it
Spatial average has <i>more</i> narrow nulls than the single positions	Vector average was used across positions	3d — it must be RMS average
Bass ~6 dB over-cut below 80 Hz	a measured simultaneous L+R sweep was used as the divisor without normalising	3c — subtract 6.0206 dB
Common and per-channel filters disagree at the 80 Hz splice	the families contain different positions, or L was equalised against R	3b and 3d
FDW appears to do nothing on a multi-position build	it was applied to the averages instead of the captures	step 2 — window all ten captures, then average

R11. Glossary

FFT bin. A measured response is a list of values at evenly spaced frequencies. The spacing here is $48000 \div 131072 = \mathbf{0.3662\ Hz}$. Nothing narrower than one bin can exist in the data, so a “2-bin feature” is 0.73 Hz wide — the finest structure the format can represent, and a guarantee that no averaging of any kind was applied.

Kernel. The weighting function inside an integral operator — the shape of the weights used when averaging neighbouring bins. For a time window, the kernel is literally the FFT of the window. Its width is what “1/6 octave” specifies. From German *Kern*, “core”. A kernel narrower than one bin changes nothing.

Minimum phase. A system whose phase is the Hilbert transform of the log of its magnitude — the unique causal system with a given magnitude and the least possible delay. Invertible with a causal, stable filter, which is why the whole procedure passes through it.

Excess phase. What is left after factoring out the minimum-phase part: magnitude exactly 1, pure delay and rotation. Contains the propagation delay, the room’s reflections, and the crossover all-pass.

Q and decay. A resonance of quality factor Q at f_0 decays 60 dB in $T_{60} = 2.2 \cdot Q / f_0$, and 40 dB in $t_{40} \approx 1.47 \cdot Q / f_0$. Inverted: to keep ringing under T milliseconds you need $Q < T \cdot f_0 / 1466$.

target	at 50 Hz	at 80 Hz	at 120 Hz
ring < 100 ms	$Q < 3.4$	$Q < 5.5$	$Q < 8.2$
ring < 50 ms	$Q < 1.7$	$Q < 2.7$	$Q < 4.1$

Note that Q is not a setting anywhere in this procedure. In the parametric / Auto EQ path it is an explicit control. In the inversion path there is no filter list and no Q field — the filter is a spectrum, and Q is an *output*, determined entirely by how narrow the features in your correction curve are. You control it only through the FDW. **The Q rule is an acceptance test, not a setting.**

Gated-tone tail. Play a steady sine, stop it abruptly, measure how long the output takes to fall 40 dB. The direct numerical version of “the music stops and the woofer keeps moving”. Unlike band-filtered decay it needs no analysis bandpass, so it has no artifact floor to subtract.

Schroeder frequency. $2000 \cdot \sqrt{(T60/V)}$ — above it the room is diffuse and statistical, below it modal and position-specific. Here, ≈ 166 Hz. It is the frequency above which a single-point measurement starts to mean something.

Part V — Automating the procedure

Everything in Parts I–IV describes doing the eleven steps by hand, in the REW GUI. `tools/rew-pipeline/rew_pipeline.py` (in this documentation/tooling repository) does the same eleven steps by driving REW’s own REST API — the same commands a person would trigger from a menu, issued over HTTP instead of a click. It does not reimplement the FDW window, the minimum-phase transform, or the divide; it asks the running REW instance to do each of those and reads back the result. That distinction is the whole point of building it this way instead of reconstructing REW’s numerics from first principles: an independent from-scratch reconstruction was tried first (`build_filters.py`, `rew_fdw_probe.py` in the same checkout) and could not be validated against REW’s actual output. Driving the real thing sidesteps the question entirely.

This part assumes you have read Parts I–IV — it explains where the script’s defaults come from and what each of its knobs corresponds to in the procedure above, not the procedure itself.

13. What it does, and does not, do for you

Requires: REW running with its API enabled (API preferences $\rightarrow$ *Start the API when REW starts*, or launched with `-api`; default port 4735), and the raw captures already loaded (imported normally, or loaded from a `.mdat` with `--session`).

Does, automatically, for one FDW cycle count per run:

- step 2 — enables the FDW at the chosen cycle count on every raw capture;
- step 3 (3c, 3d) — the mono sum at each position, then the three spatial RMS averages;
- step 3e — bakes the crossover correction (X801) into the channel averages;
- step 4 — minimum phase, first time, cal included;
- step 5 — builds the target, if one is not already loaded under the configured name (§16);
- step 6a — checks the minimum-phase copies preserved |H|, and warns if not (§19 has the one case this actually caught);
- step 7, 8, 9 — the division, minimum phase a second time (cal excluded), bakes the crossover in last;
- step 10 — trims to set latency, exports the filter WAVs and the text responses open-media-drc’s deployment tooling needs (§17);
- step 11 — runs `tools/drc_acceptance.py` on the result and reports pass/fail.

Does not do for you, on purpose:

- **3a, 3b** — collapsing repeat sweeps and (optional) level alignment. Do these once, by hand, before the captures the script consumes exist; they are judgment calls about your own measurement set, not something to automate per run.
- **Choosing a target shape or house curve.** §16 covers what the script sets on its own, but the actual curve — flat, a house curve, a deliberate scoop — is a decision, not a default to compute. Build one by hand in REW once and point `--target-title/target_title` at it to have every run reuse it, or let the script build the plain default described in §16.
- **step 11’s actual accept/reject judgment.** The script runs `tools/drc_acceptance.py` and reports its verdict; it does not loop over FDW cycles or LF-tail slopes looking for a pass. That decision — and the re-measurement at more than one position the guide’s step 11 also asks for — stays with you.

14. Input measurements, and the config file

Naming convention, overridable everywhere below: a centre pair L0/R0, optionally 1–4 more position pairs L1/R1 ... L4/R4, each *n* cm off centre (§6). If a real simultaneous sweep L0+R0 exists it is used as the centre mono sum (3c), normalised -6.0206 dB automatically; otherwise the centre sum is L0 vector-averaged with R0, same as every other position.

Config file. Every DSP flag lives in one TOML file instead of a long command line: copy `../DRC-doc/tools/rew-pipeline/rew_pipeline.example.toml` into the project as `rew_pipeline.toml` (loaded automatically if present) or pass `--config some-file.toml`. A command-line flag always overrides the file; a key simply absent from the file falls back to the script's own built-in default — there is no separate “not set” value to get wrong. This project's own `rew_pipeline.toml` carries its actual measurement titles (L.0/R.0/L 120.green.{n}, not the generic L0/L{n}), so, when run from that project directory, a plain `../DRC-doc/tools/rew-pipeline/rew_pipeline.py --tag fdw8 --output output/fdw8` is a complete, correct run with no other flags.

.toml table	key	script flag	meaning	default
[measurements]	center_l, center_r	--center-l, --center-r	the centre pair's titles	L0, R0
	pos_l_pattern, pos_r_pattern	--pos-l-pattern, --pos-r-pattern	{n}-templated titles for the 4 off-centre positions	L{n}, R{n}
	num_positions	--num-positions	how many of those 4 to use, 0–4	4
	sum_c_title	--sum-c-title	a real simultaneous L0+R0 sweep, if you have one	unset → vector-average
	x801_title, x801_wav	--x801-title, --x801-wav	the crossover correction, reused if already loaded, else imported	X801, X801.wav
[fdw]	cycles	--fdw-cycles	step 2's FDW width	8.0
[target]	target_title	--target-title	reused as-is if already loaded; else built (§16) under this name	Target LR.RMS.AVG
	lf_cutoff_hz, lf_slope_db_per_oct	--target-lf-cutoff, --target-lf-slope	step 5's target LF cutoff	10 Hz, 24 dB/oct
	house_curve	--house-curve	a house-curve file (freq/dB pairs, same format as <code>../DRC-doc/tools/housecurve/house-curve-*.txt</code>), loaded before the target is built	unset → flat
	house_curve_log_interpolation	--house-curve-log-interpolation	REW's own flag, set explicitly since there is no way to read back a prior value	true
[minphase1]	lf_tail_corner_hz, lf_tail_slope_db_per_oct	--lf1-corner, --lf1-slope	step 4's LF tail	16 Hz, 12 dB/oct
	hf_tail_corner_hz, hf_tail_slope_db_per_oct, hf_tail_frequency_warping	--hf1-corner, --hf1-slope, --hf1-warping	step 4's HF tail	unset → off
	same four keys	--lf2-*, --hf2-*	step 8's tails	16 Hz @ 0 dB/oct, HF off
	replicate_data	--replicate-data	REW's flag for whichever tail is disabled	false
[division]	common_low_hz, common_split_hz, upper_hz	--common-low, --common-split, --upper	step 7's three band edges	25, 80, 225 Hz

Cal file effects (included in minimum phase #1, excluded in #2, R3) and max gain (0.0 dB, cut-only, throughout §11) are not configurable — the guide's reasoning for both is a property of the procedure, not a per-run choice.

Optional session cleanup

Normal runs reuse an already loaded input session and replace their own same-tag results. Two additional CLI-only options allow a clean rebuild:

- `--replace-session` requires `--session PATH`: it deletes all currently loaded REW measurements before loading that session.
- `--raw-only-session` keeps only the configured raw centre and off-centre L/R captures, plus the simultaneous centre sum if configured. It removes prior targets, filters, arithmetic results and imported X801; the build recreates these from the captures and configured X801 WAV. Locked traces that REW refuses to delete are reported and excluded from the build. Missing required raw inputs cause an error.

Both options are off by default. They modify the loaded REW session, so save any work you need before using them. These explicit cleanup options are broader than normal same-tag cleanup.

15. Steps 2–10, one API call each

The mapping is close to literal — each guide step is one or two REST calls per trace, not a reimplementaion of what they compute:

guide step	REST call	REW command / endpoint
2 — FDW	POST /measurements/:id/ir-windows	{"addFDW": true, "fdwWidthCycles": N}
3c — mono sum	POST /measurements/process-measurements	"processName": "Vector average"
3d — spatial average	same endpoint	"processName": "RMS average"
3e, 7d, 9 — multiply	same endpoint	"processName": "Arithmetic", "function": "A * B"
4, 8 — minimum phase	POST /measurements/:id/command	"command": "Minimum phase version"
5 — target	POST /measurements/:id/eq/command	"Calculate target level", then "Generate target measurement"
7 — division	POST /measurements/process-measurements	"processName": "Arithmetic", "function": "A / B", maxGain/lowerLimit/upperLimit
10 — trim	POST /measurements/:id/command	"command": "Trim IR to windows"
— export	GET /measurements/:id/impulse-response	?windowed=true&normalised=false&unit=percent, written to WAV directly

Every one of these that produces a new trace does so under a fresh measurement, named `<role>.<tag>` (LX.fdw8, LFilter.fdw12, ...) so that two runs with different `--tags` coexist in the same REW session without clashing, and a rerun of the *same* tag replaces only its own prior measurements — never the raw captures, the target, or X801.

Optional refinement PEQs: 82 Hz and 530 Hz

The refinement extension adds two configurable, common L/R peaking cuts to the inversion. Relative to DRC-120.green commit f4fddf9, this is new behavior in the local pipeline changes documented on 2026-09-13; that committed version does not offer these options. The underlying inversion bands, target construction and FDW defaults are unchanged.

Filter	Frequency	Gain	Q	Channels
Remaining bass peak	82 Hz	−1.0 dB	2.0	Both, identically
Broad upper-bass/lower-midrange peak	530 Hz	−1.5 dB	2.0	Both, identically

Off by default. Enable with `--refinement-peq`, or put this in the active TOML config (`rew_pipeline.example.toml` documents the same settings with `enabled = false`):

```
[experiments]
enabled = true
peq82_frequency_hz = 82.0
peq82_gain_db = -1.0
peq82_q = 2.0
peq530_frequency_hz = 530.0
peq530_gain_db = -1.5
peq530_q = 2.0
```

The corresponding CLI controls are `--peq82-frequency`, `--peq82-gain`, `--peq82-q`, `--peq530-frequency`, `--peq530-gain`, and `--peq530-q`. Gains must be non-positive and Q must be positive. There is one enable switch for both PEQs; setting one gain to 0 makes that cut neutral when exporting a single-cut trial. There is currently no `--no-refinement-peq` flag: if TOML enables the feature, set `enabled = false` there or select a config that leaves it disabled to obtain a baseline run. Supplying PEQ parameters alone does not enable the feature.

When enabled, both cuts are baked into the exported filters. After forming $FL = F.common \times Fper_L$ and $FR = F.common \times Fper_R$, the pipeline constructs explicit REW PK filter responses and multiplies their combined response into both channels:

```
FL.exp = FL × PEQ.both
FR.exp = FR × PEQ.both
→ minimum phase #2 → multiply by X801 → trim → export
```

Thus the normal FLX-trimmed-48k.wav / FRX-trimmed-48k.wav, filter text exports and L.filtered / R.filtered / LR.filtered predictions include both cuts. No additional playback PEQ is required. This added stage also applies the 530 Hz cut even though the inversion’s upper band edge remains 225 Hz. It does not change the inversion target or extend the division band.

The supporting `rew_client.py` methods replace the EQ rows on the working filter measurements, clear unused rows, read back the requested settings, and invoke REW’s **Generate filters measurement** command. The response is then explicitly cascaded into the correction; it is not merely an EQ-panel setting attached to a plot.

For inspection, the enabled run retains these tagged trace families in REW:

- PEQ.82, PEQ.530, PEQ.both: standalone EQ responses.
- LFilter.baseline / RFilter.baseline, FLX.baseline / FRX.baseline, and FLX-trimmed.baseline / FRX-trimmed.baseline: the correction rebuilt through the same minimum-phase, X801 and trimming steps without either cut.
- L.filtered.baseline, R.filtered.baseline, LR.filtered.baseline: baseline predictions.
- L.filtered.eq82, R.filtered.eq82, LR.filtered.eq82, and the matching .eq530 family: individual-cut predictions, made by multiplying the finished baseline predictions by each standalone PEQ. These are comparison traces, not separately exported single-cut filter builds.
- The ordinary L.filtered, R.filtered, LR.filtered: the combined candidate used for export.

The manifest records the parameters and comparison UUIDs; `--save-mdat` preserves the session for later inspection. To make room for comparisons, the enabled run deletes its temporary per-position sum and LX/RX construction traces after their results have been incorporated. Same-tag cleanup also recognises the new PEQ and comparison trace names.

For example, from the project checkout, build the FDW6 combined candidate:

```
../DRC-doc/tools/rew-pipeline/rew_pipeline.py --fdw-cycles 6 --refinement-peq \
--tag clean-fdw6-peq --output output/clean-fdw6-peq --save-mdat
```

The existing `output/clean-fdw6-peq/manifest.json` records FDW 6 and both cuts at the values above. The active `DRC-120.green/rew_pipeline.toml` at this writing has no `[experiments]` section and still selects FDW 8, so a normal run enables neither cut. Predictions are calculated from the input captures; judging the tonal change still requires listening and acoustic verification of the final filters.

16. Building the target

Answering directly, since it is the one step with real judgment in it:

- **House curve:** none, by default — the built target is flat (step 5’s “no scoop” option), because a curve is a decision this script should not make silently. Pass `--house-curve/house_curve` to load one first (REW’s `/eq/house-curve` endpoint) — `../DRC-doc/tools/housecurve/house-curve-harman-fuller.txt` is one such file, see step 5’s “four candidate shapes”.
- **Target level:** yes — “Calculate target level”, exactly step 5’s “press Calculate and let REW set it”, run on the RMS average of LX/RX (the same trace step 5 uses).
- **Then “Generate target measurement”:** yes, immediately after, on the same trace — this is the call that actually produces the target measurement the division in step 7 consumes.
- **Target type / shape:** “Full range”, always, not “Flat” (rejected — REW’s valid shapes are Full range, Bass limited, Subwoofer, Driver, None) and not REW’s own default of “Subwoofer”. That default matters: it silently attaches a bass-management crossover (REW’s `TargetSettings.bassManagementCutoffHz/bassManagementSlopedBPerOctave`, visible only once you inspect the object — see R12 below), and a target built under it produced a visibly worse filter — sharpest feature Q 12–13 against a $Q \leq 12$ gate, and +11...+24 ms of group delay where the intended build shows single-digit ms — until this was found and fixed. Always Full range for a two-channel main-speaker target.
- **LF cutoff / slope:** 10 Hz at 24 dB/oct by default (`target_lf_cutoff / target_lf_slope`), matching step 5’s “move it to 5–10 Hz, below the correction band” — REW’s own factory default is 20 Hz, which sits inside the match range and bends the target there (step 5 has the measured table).

If a target under the configured name is already loaded — for instance one you shaped by hand in REW with a deliberate scoop — none of the above runs; the script reuses it as-is. That is the intended way to try a non-default target shape.

17. Output files

Exactly the ten files `open-media-drc`'s `scripts/new_filter_design.py` resolves by name (its own `FILTERS_` AND `DRC.md` and the `TXT_NAMES` / `AGGREGATE_NAMES` / `WAV_PATTERNS` tables in that script), written into `--output/output`, nothing else:

<code>FLX-trimmed-48k.wav</code>	<code>FRX-trimmed-48k.wav</code>	the two filters -- BruteFIR input	
<code>FLX-trimmed.txt</code>	<code>FRX-trimmed.txt</code>	their frequency response	
<code>L.txt</code>	<code>R.txt</code>	<code>LR.txt</code>	measured, before correction
<code>L.filtered.txt</code>	<code>R.filtered.txt</code>	<code>LR.filtered.txt</code>	raw (no-FDW) capture x filter

plus `manifest.json` (every parameter and measurement UUID from the run) and `acceptance.txt` (`tools/drc_acceptance.py`'s own output) — both ignored by `new_filter_design.py`, which matches files by name and does not mind extras. Deliberately **not** exported: a text trace for every intermediate measurement the pipeline builds (`LX`, `LX-MP`, `F.common`, ...) — those stay in REW for inspection if you want them, but going to a design bundle would just be clutter next to the eight curves anything downstream actually plots.

This directory is a complete, ready-to-import export set — the same shape as a `.txts` folder exported by hand from REW, beside the `.mdat` it came from. Point `new_filter_design.py` (see its own `--help` and `../open-media-drc/scripts/README.md`) at it to deploy; `--dry-run` runs every check, including the filter `TXT/WAV` residual, without writing anything. A real (non-dry-run) deployment additionally wants the ten files and the `.mdat` committed to git and `--mdat <session>.mdat` naming the session, both left to you rather than assumed by `rew_pipeline.py`.

18. Rebuilding at a different FDW, or after a tweak

```
../DRC-doc/tools/rew-pipeline/rew_pipeline.py --fdw-cycles 12 --tag fdw12 --output output/fdw12
```

is the whole of it — a new `--tag` keeps the new run's measurements alongside the old ones in REW rather than overwriting them, so both can be inspected side by side before choosing. Rerunning the *same* tag (after editing `rew_pipeline.toml`, say, to try a different LF-tail slope) replaces that run's own measurements in place.

19. Quirks the API did not document

Found by testing against this project's own reference build, not written down anywhere in the API help:

- **Measurement index numbers shift on every add or delete** (the API docs do say indexing is "NOT recommended" but not why it actually bites): a value captured as an index and used several steps later, after other measurements have been created or removed, can silently refer to the wrong trace. Every reference the script keeps across steps is the measurement's **UUID**, never its index.
- **Response copy does not carry the source's notes** — only its own "Copy of <title>" line — so the geometry comments a downstream tool might look for (front-wall/speaker distance, marker colour — §17, R12) have to be copied over explicitly afterward.
- **Response copy does not detach the IR-window/FDW state either.** Taking a "no-FDW" snapshot of a capture *before* enabling FDW on the original, expecting the copy to be unaffected by what happens to the original next, produced a copy that showed FDW enabled anyway once the original was mutated. The fix matches what actually happened in the reference session's own history: build the *entire* chain first with FDW on throughout, and only at the very end — once nothing downstream still needs it — turn FDW off on the original captures and take the snapshot then.
- **/measurements/:id/frequency-response could not be made to return genuinely unsmoothed data** in testing, regardless of a `smoothing=None` query parameter or a "Smooth" command with "None" issued first — it came back at whatever log-spaced default (1/48 octave, 96 points/octave) the measurement last displayed at in the GUI. Every text export here is therefore computed directly from `/measurements/:id/impulse-response` instead (a plain FFT of the full-length impulse is unsmoothed and linear-spaced by construction), not from that endpoint.
- **A filter's exported text and its WAV must agree on where t=0 is**, and it is not the buffer start. `new_filter_design.py`'s own residual check only searches for an alignment within 16 samples of the WAV's peak; a text export phase-referenced to the buffer's sample 0 instead of to that peak sample produced a spurious ~104° RMS phase residual — appearing as a data-quality failure — even though the magnitude matched to 0.0002 dB. Referencing phase to the peak sample fixed it to exactly 0.

R12. Appendix — how the REW API was explored

The API is self-documenting to a useful degree, and every non-obvious setting used above was found this way rather than guessed. This section is the method, so it is repeatable for whatever the next unfamiliar endpoint turns out to be.

Start with the specification. GET /doc.json on a running instance returns the full OpenAPI/Swagger document — every path, every request/ response schema, by name. For a data model with fields whose meaning is not obvious from the prose help at GET / (rendered from the same source as localhost:4735's own browsable page), the schema in doc.json is the fastest way to see what a POST body can actually contain, for instance:

```
"TargetSettings": {
  "type": "object",
  "properties": {
    "shape": {"type": "string"},
    "bassManagementSlopedBPerOctave": {"type": "integer"},
    "bassManagementCutoffHz": {"type": "integer"},
    "lowFreqSlopedBPerOctave": {"type": "integer"},
    "lowFreqCutoffHz": {"type": "integer"},
    "lowPassCrossoverType": {"type": "string"},
    "highPassCrossoverType": {"type": "string"},
    "lowPassCutoffHz": {"type": "integer"},
    "highPassCutoffHz": {"type": "integer"}
  }
}
```

That is where **bassManagementCutoffHz** and its neighbours were first seen — not from any narrative documentation, which never mentions them by name — by reading the schema for the object GET /measurements/:id/target-settings returns. The type only says the field exists and that it is an integer; it does not say which shape values are valid, or that bassManagementCutoffHz only matters for some of them.

Read a live object back before changing it. GET /measurements/:id/target-settings on an ordinary loaded measurement returned:

```
{
  "shape": "Subwoofer",
  "bassManagementSlopedBPerOctave": 12,
  "bassManagementCutoffHz": 80,
  "lowFreqSlopedBPerOctave": 24,
  "lowFreqCutoffHz": 10,
  "lowPassCrossoverType": "L-R2",
  "highPassCrossoverType": "L-R2",
  "lowPassCutoffHz": 1000,
  "highPassCutoffHz": 100
}
```

"Subwoofer" as the factory-default shape, for a full-range two-channel target, was the finding that mattered — not something to have assumed, since nothing in the prose documentation says what REW's default target shape is for a fresh measurement. This is also what made the actual bug visible: building a target under this default and running tools/drc_acceptance.py against the resulting filter showed a materially worse result (§16) than the existing reference build, which pointed straight back at this object.

Let the API's own validation enumerate the choices. The schema says shape is a string; it does not say which strings. Rather than search the GUI for the exact wording, POST a plausible guess and read the error:

```
$ curl -X POST localhost:4735/measurements/<uuid>/target-settings \
  -d '{"shape": "Flat"}'
{
  "message": "Flat is not a valid target shape",
  "validValues": ["Full range", "Bass limited", "Subwoofer", "Driver", "None"]
}
```

REW's API consistently does this: an invalid enum value comes back with the complete valid list, so a wrong guess is one request, not a search through the GUI or the help pages. The same pattern found the valid Arithmetic "function" strings (GET /measurements/arithmetic-functions) and the valid smoothing values

(GET /measurements/frequency-response/smoothing-choices) before any POST was needed at all — some endpoints hand you the enumeration directly, without needing to trigger a validation error first.

The FDW is a field on the IR-window object, not a separate endpoint. GET /measurements/:id/ir-windows on a capture that already had it enabled (one from the reference session, inspected specifically to confirm this) returned:

```
{
  "leftWindowType": "Rectangular", "leftWindowWidthms": 500,
  "rightWindowType": "Rectangular", "rightWindowWidthms": 1000,
  "refTimems": 0,
  "addFDW": true, "fdwWidthCycles": 8.0
}
```

— confirming both that addFDW/fdwWidthCycles are the only two fields that matter for step 2, and, independently, that this project’s own reference build really did use exactly 8 cycles and the rectangular 500/1000 ms windows Part III’s step 2 assumes throughout. Setting the FDW is then one POST to the same endpoint with just those two fields — the API explicitly documents partial objects being accepted by POST, so the rectangular window fields do not need to be repeated.

Minimum phase’s tail parameters came from the prose help, verified against a real object. Unlike the two above, /measurements/:id/command with "Minimum phase version" *is* documented with a worked example in the HTML help (analysis/rew-api-help.txt in the project checkout, saved from GET /), naming "include cal", "append lf tail", "lf tail start", "lf tail slope" explicitly. What the help does not give is which numbers this project’s own reference filter actually used — for that, 120.green.multipoint.FDW8.mdat’s own measurement *notes* were the source, not the API: each -MP and Filter measurement in that session carries a REW-generated note stating its own tail settings verbatim, e.g.

```
Minimum phase copy of LX
Cal file effects included
LF tail from 16 Hz at 12 dB/octave
No HF tail
```

read via GET /measurements/:id on that specific reference measurement. Every “16 Hz @ 12 dB/oct” / “16 Hz @ 0 dB/oct” default in §14’s table is transcribed from notes like this one, not chosen or guessed — the API only had to be asked to *set* the same values back, on a fresh build, once they were known.

When neither the schema nor an error message settles it, test against a known-good file. The one genuinely open question — what phase convention a filter’s text export must use to satisfy new_filter_design.py’s own TXT-vs-WAV residual check (§19’s last item) — was not answered by any REW endpoint at all, because the check lives in the *other* project. It was settled by importing deploy_filter.filter_spectrum and response_metrics directly and running them against a real WAV/TXT pair at a few candidate delays:

```
for delay in [0, 8192, -8192]:
    ...
    # delay=0          rms_phase_deg = 0.0015    (this script's own convention)
    # delay=8192 (peak) rms_phase_deg = 0.0000    (what the checker expects)
    # delay=-8192      rms_phase_deg = 104.3
```

— a one-sample offset either side of the correct value at 8192 already produces ~104° of residual (a phase check this sensitive is doing its job: one sample at 24 kHz is half a cycle), which is what made the fix unambiguous once found.